




USC University of
Southern California

Masoud Moshref, Minlan Yu
Ramesh Govindan, Amin Vahdat

Trumpet: Timely and Precise Triggers in Data Centers



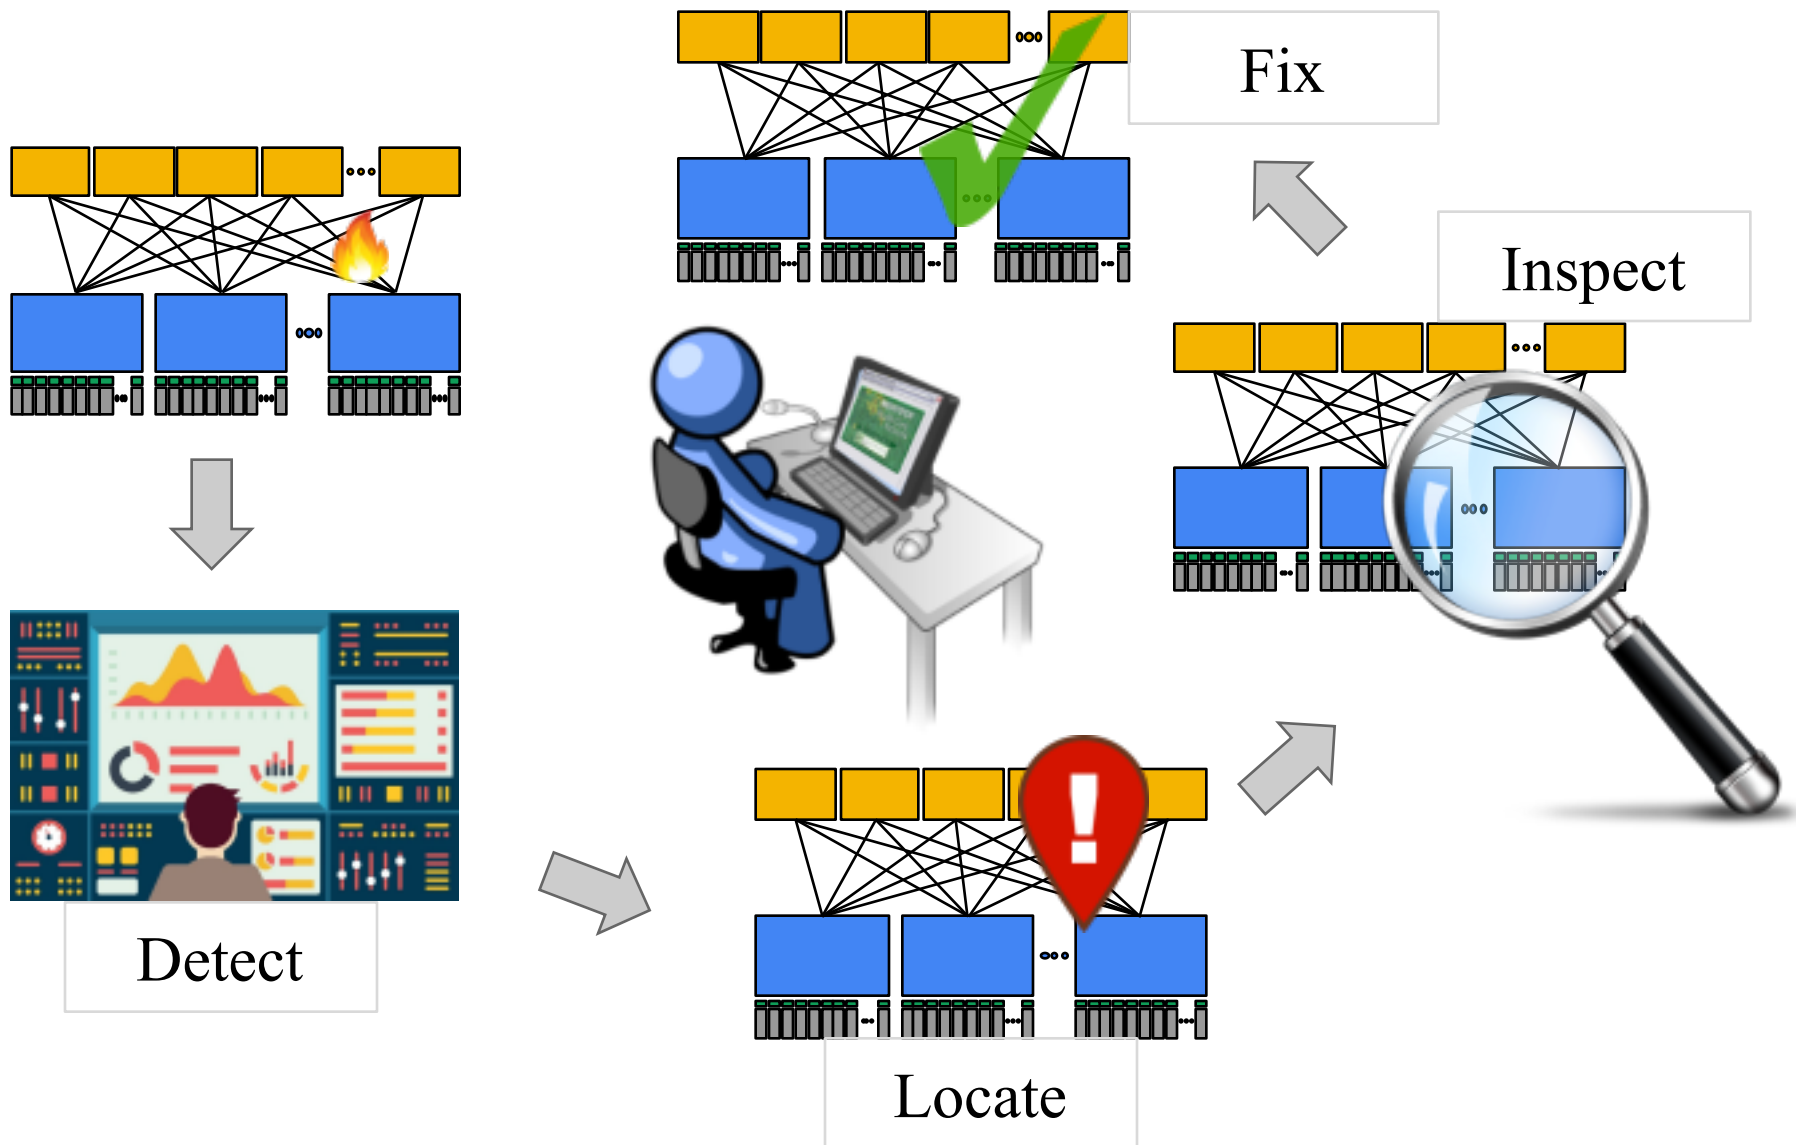
Evolve or Die, SIGCOMM 2016

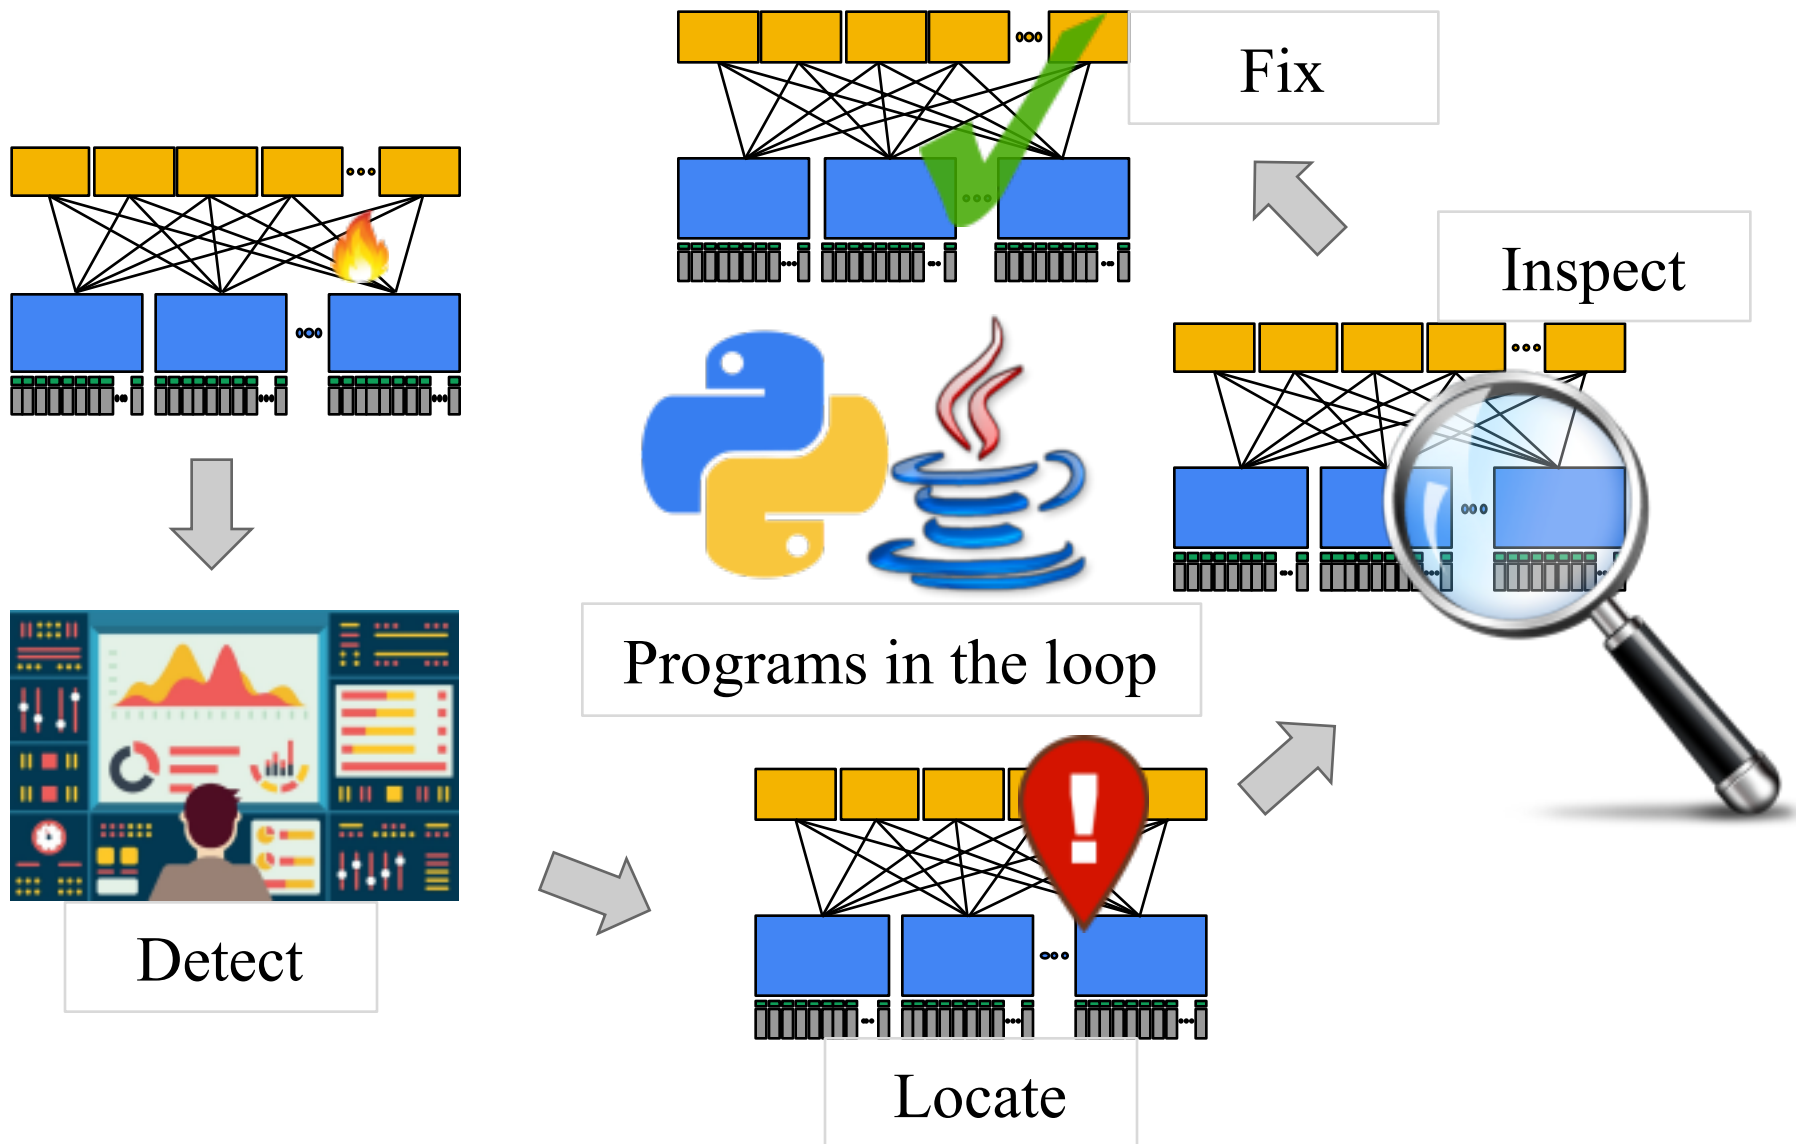
Long failure
repair times in
large networks



Human-in-the-loop failure
assessment and repair

Humans in the Loop



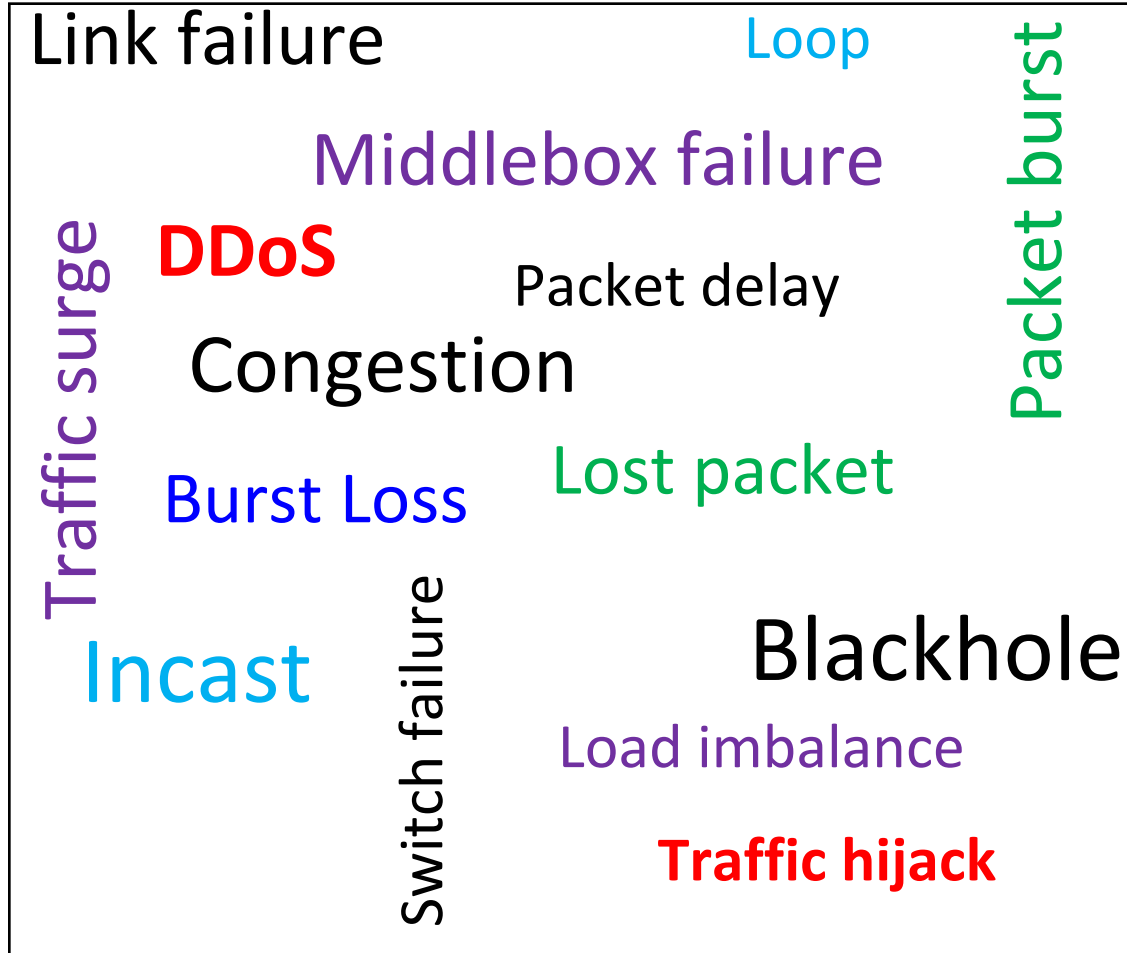




Detect



A framework for **programmed** detection
of *events* in large datacenters





- ❖ Availability
- ❖ Performance
- ❖ Security



Detect

**Aggregated, often sampled
measures of network health**

Detecting Transient Congestion




40 ms burst

Timeouts lasting
several 100 ms

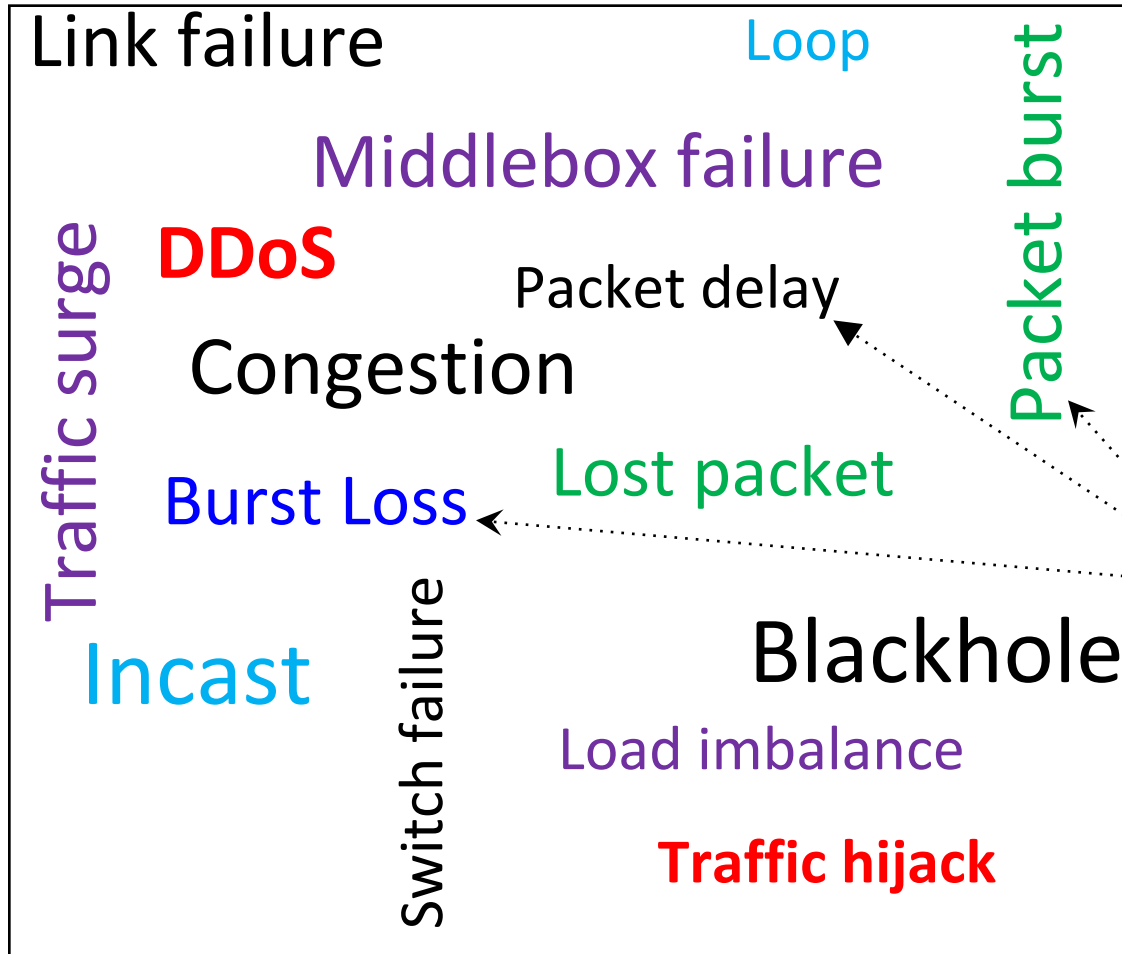
Fine
Timescale
Events

Detecting Attack Onset

Fine
Timescale
Events



Did this **tenant** see a sudden increase in traffic over the last few **milliseconds**?



Some event definitions may require **inspecting every packet**

Eventing Framework Requirements

Expressivity

- Set of possible events not known *a priori*

Fine timescale eventing

- Capture transient and onset events

Per-packet processing

- Precise event determination

Because data centers will require high availability and high utilization

Where do we place eventing functionality?

Switches



NICs



Hosts



- ❖ Are programmable
- ❖ Have processing power for fine-time scale eventing
- ❖ Already inspect every packet

*We explore the design of a
host-based eventing
framework*

Research Questions

What eventing architecture permits programmability **and** visibility?

How can we achieve precise eventing at fine timescales?

What is the *performance envelope* of such an eventing framework?

Research Questions

What eventing architecture permits programmability **and** visibility?

How can we achieve precise eventing at fine timescales?

What is the performance envelope of such an eventing framework?

Trumpet has a logically centralized *event manager* that aggregates *local events* from per-host *packet monitors*

For each packet matching **Filter**

group by **Flow-granularity**

and report every **Time-interval**

each group that satisfies **Predicate**

Flow volumes, loss rate, loss
pattern (bursts), delay

Is there any flow sourced by a service that sees a burst of losses in a small interval?

For each packet matching

Service IP Prefix

group by

5-tuple

and report every

10ms

any flow whose

$\text{sum}(\text{is_lost} \ \& \ \text{is_burst}) > 10\%$

Is there a **job** in a **cluster** that sees **abnormal traffic volumes** in a **small interval**?

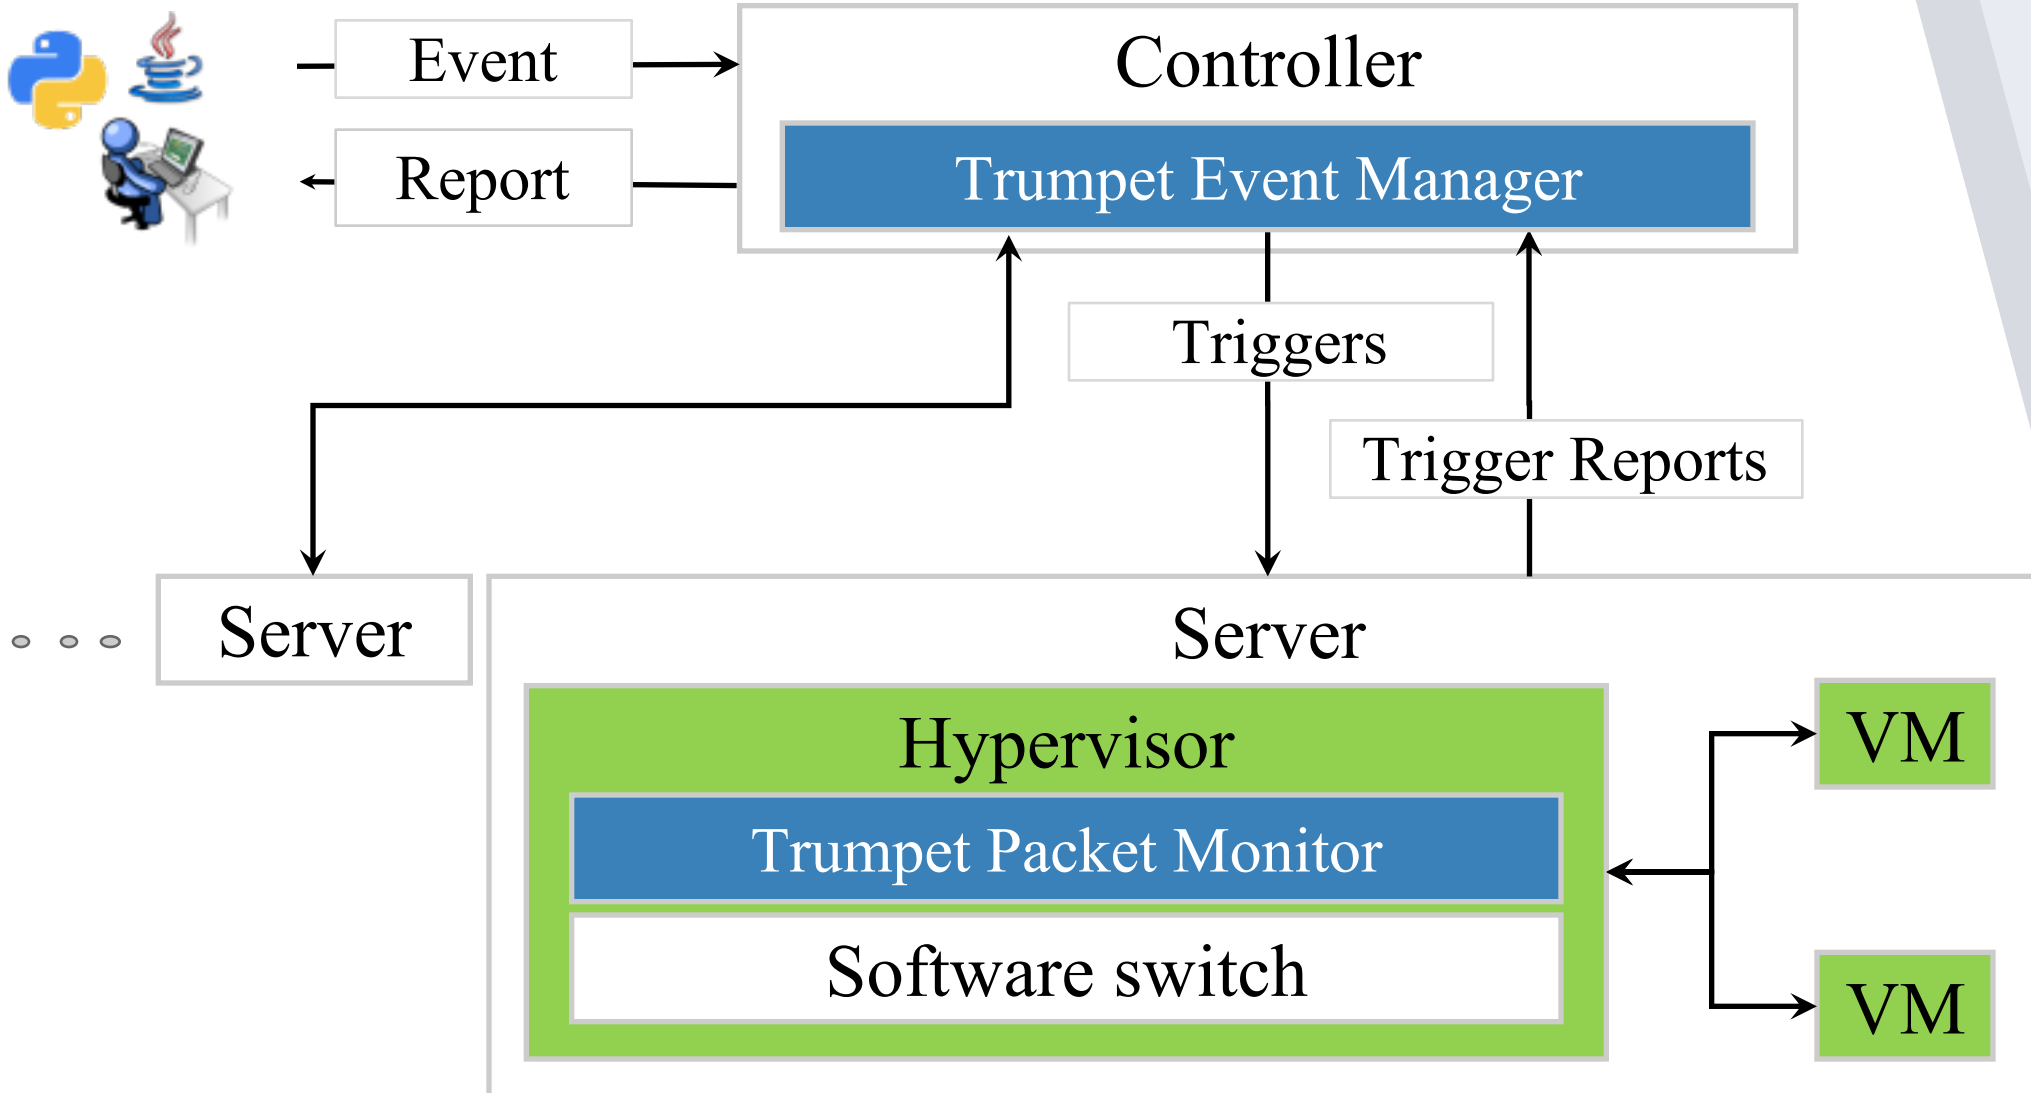
For each packet matching

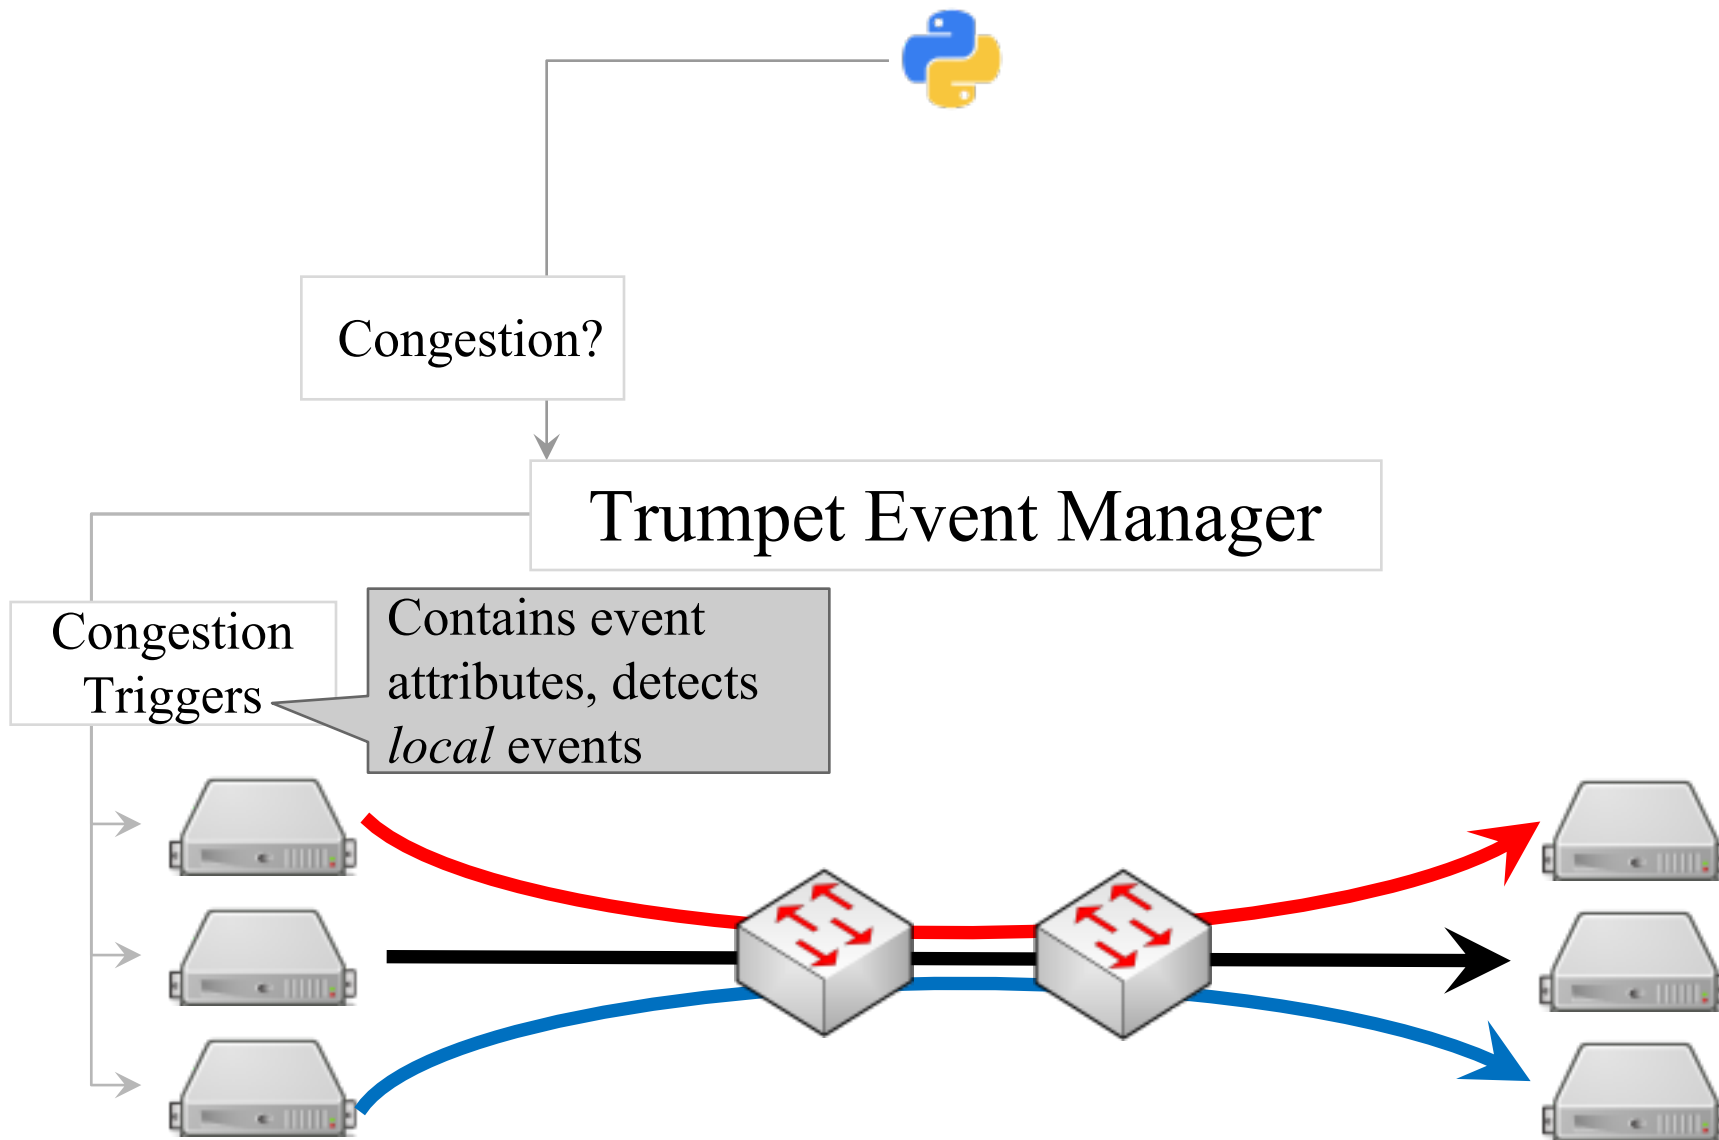
**Cluster IP Prefix
and Port**


group by **Job IP Prefix**

and report every **10ms**

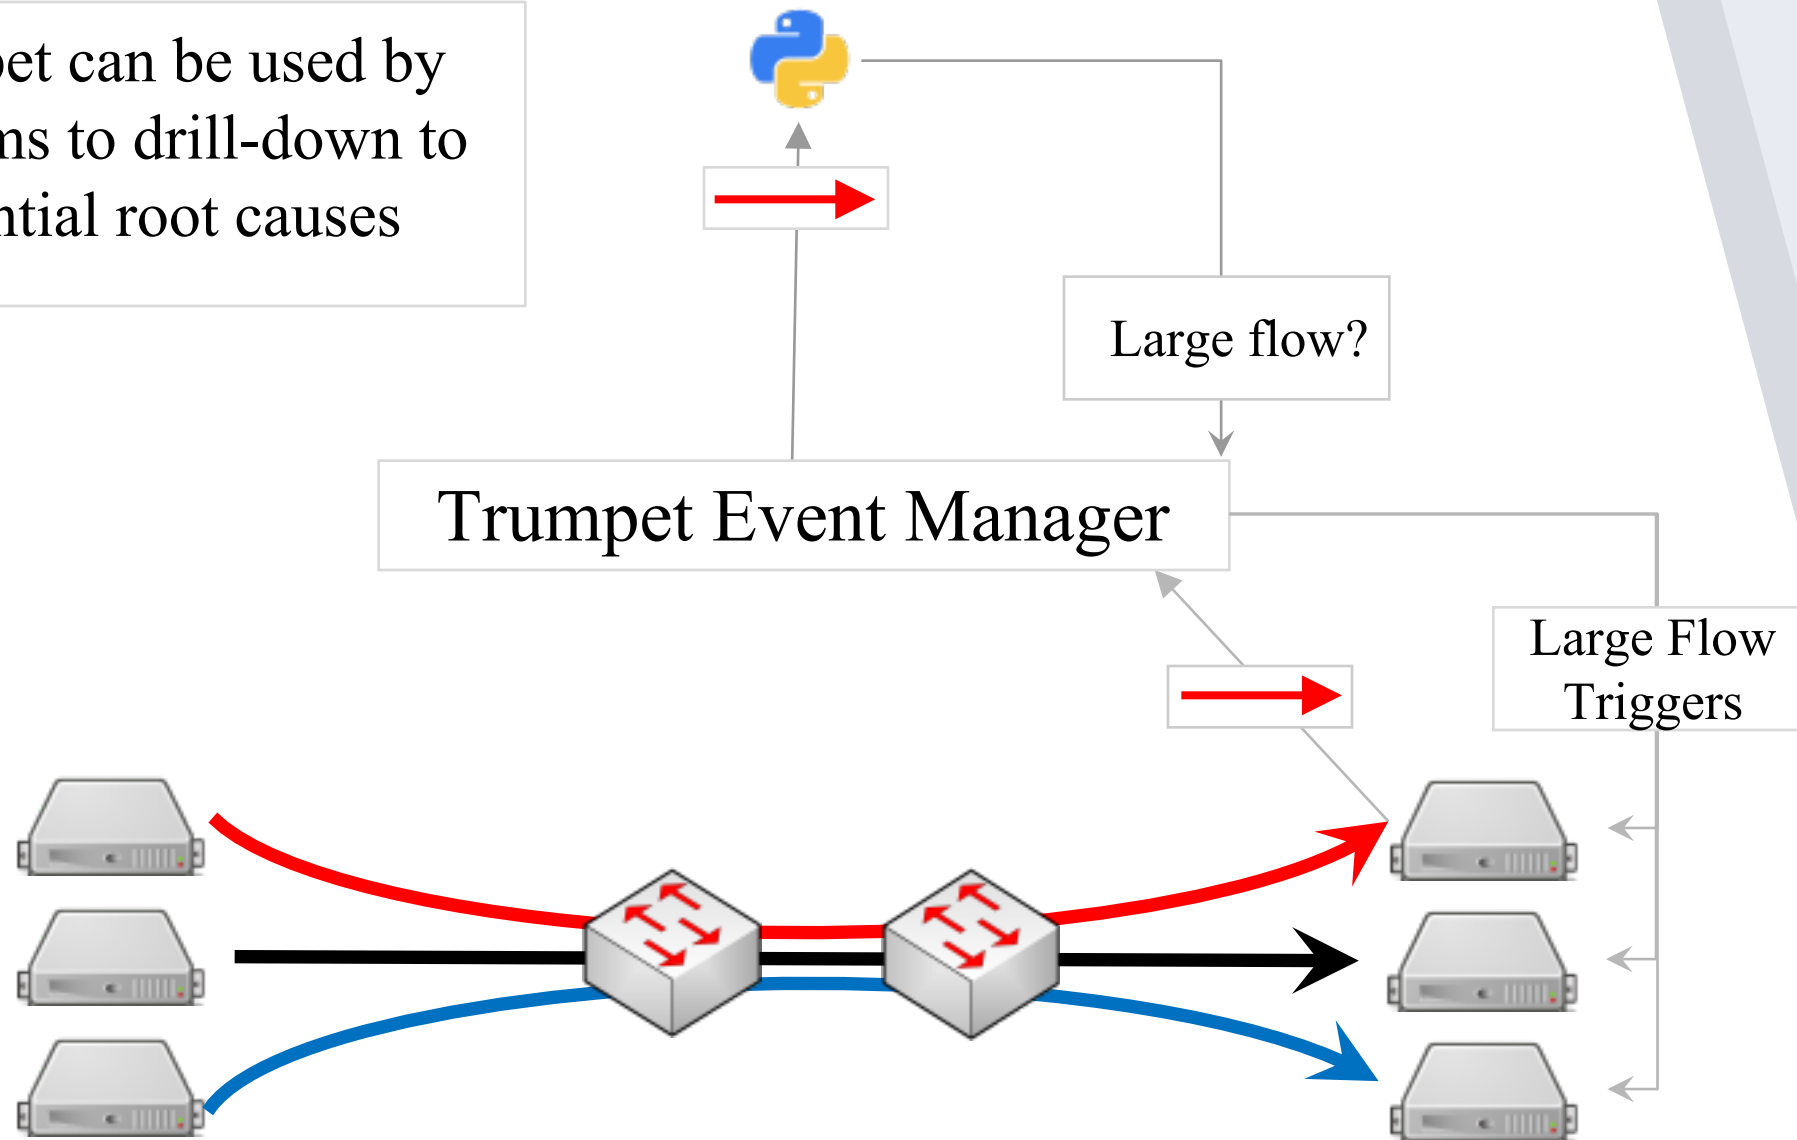
any job whose **sum (volume) > 100MB**







Trumpet can be used by programs to drill-down to potential root causes



Trumpet
Event
Manager


Research Questions

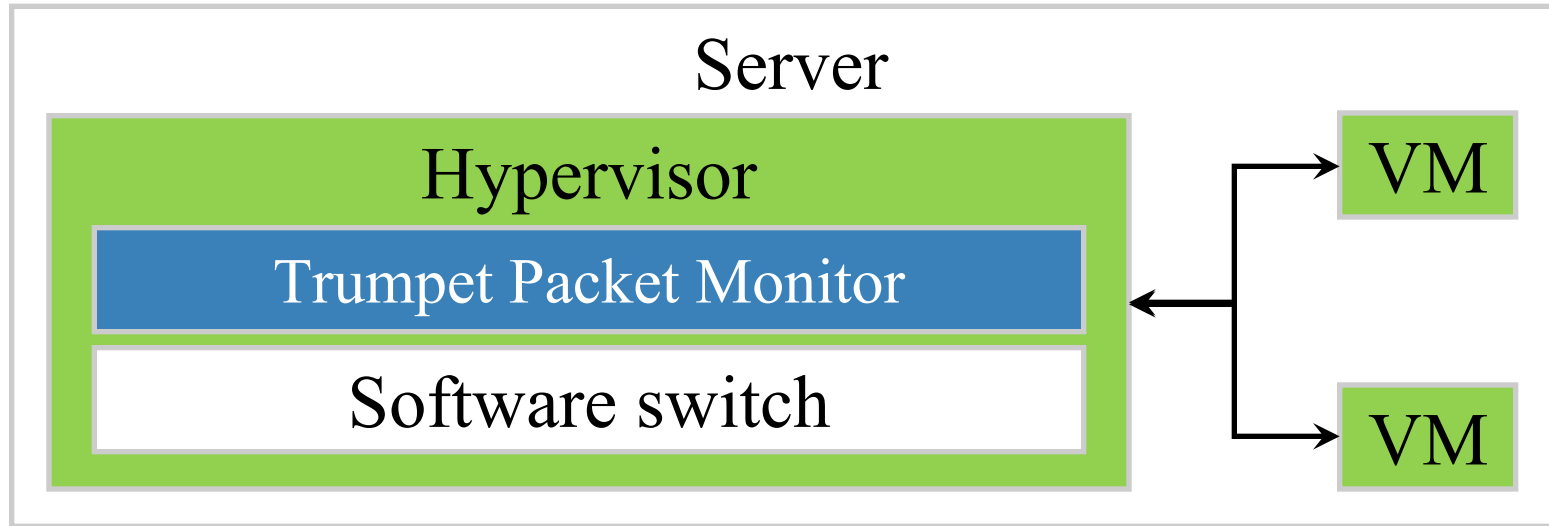
What eventing architecture permits programmability and visibility?

How can we achieve precise eventing at fine timescales?

What is the *performance envelope* of such an eventing framework?

The monitor optimizes packet processing to inspect every packet and evaluate predicates at fine timescales





Piggyback on CPU core used by software switch

- ❖ Conserves server CPU resources
- ❖ Avoids inter-core synchronization

*Can a single core monitor thousands
of triggers at full packet rate (14.8
Mpps) on a 10G NIC?*

Two Obvious Tricks

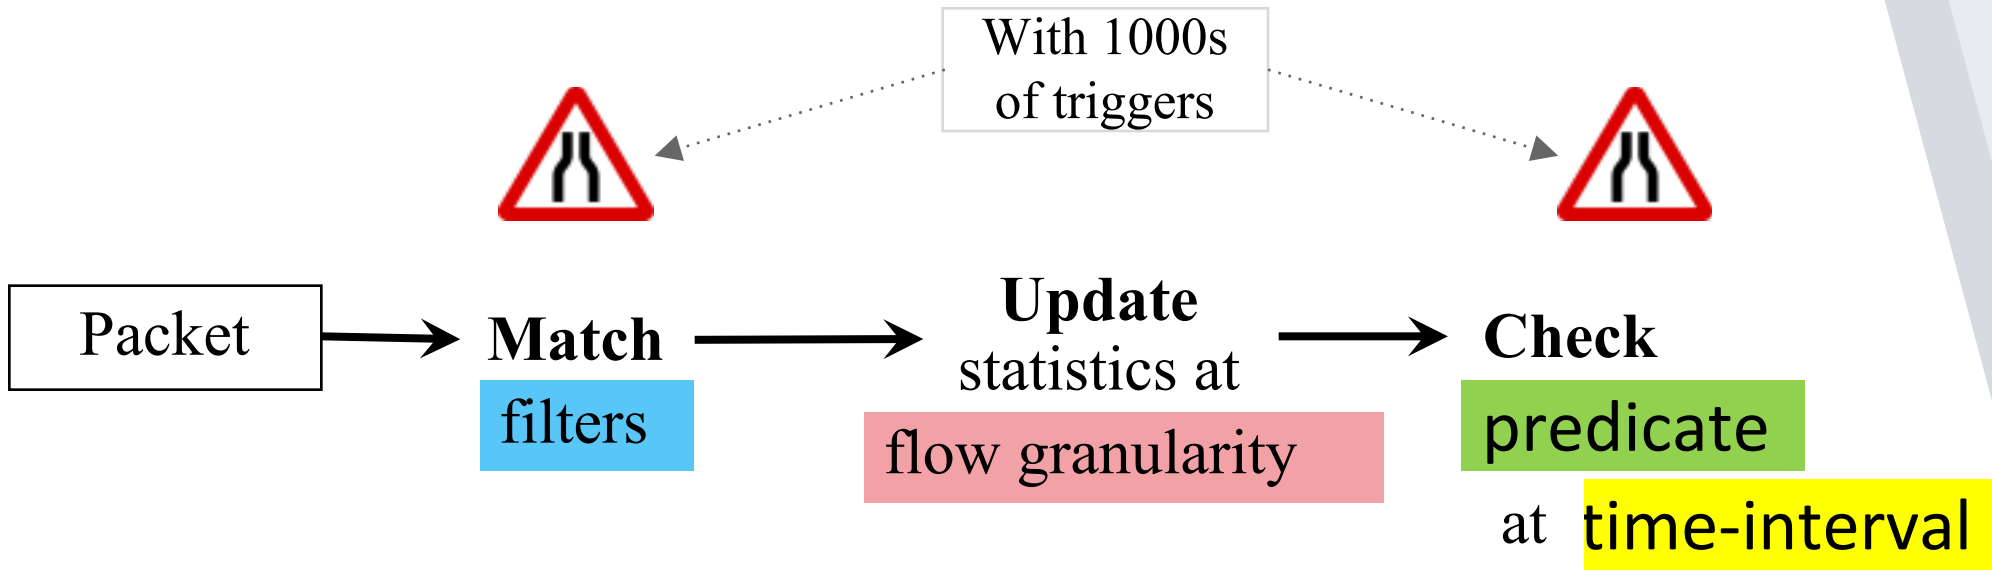
Use kernel bypass

- ▶ Avoid kernel stack overhead

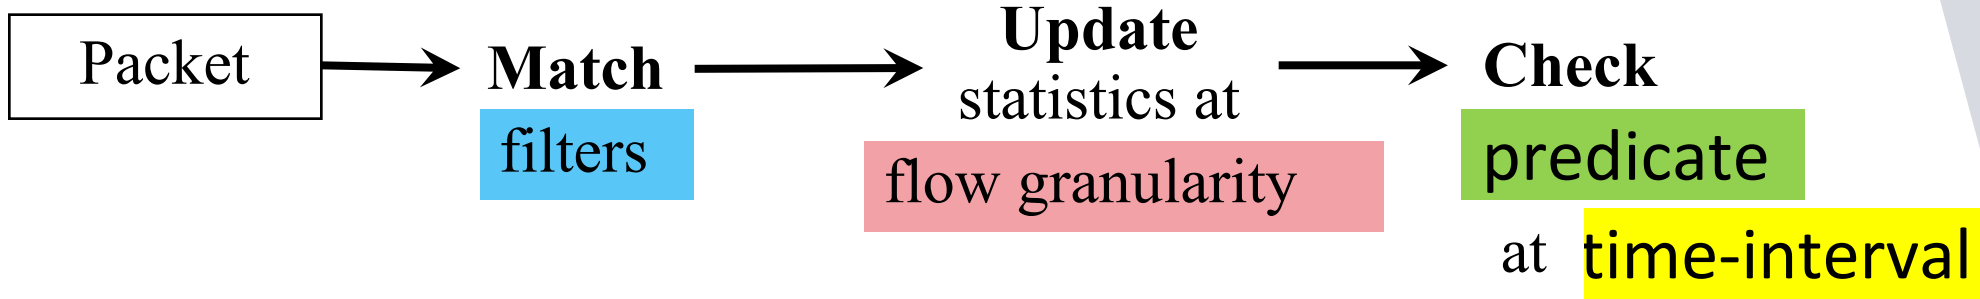
Use polling to have tighter scheduling

- ▶ Trigger time intervals at 10ms

Necessary, but far from sufficient....




Filter	Flow granularity	Time interval	Predicate
Source IP = 10.1.1.0/24	5-tuple	10ms	Sum(loss) > 10%
Source IP = 20.2.2.0/24	Service IP prefix	100ms	Sum(size) < 10MB

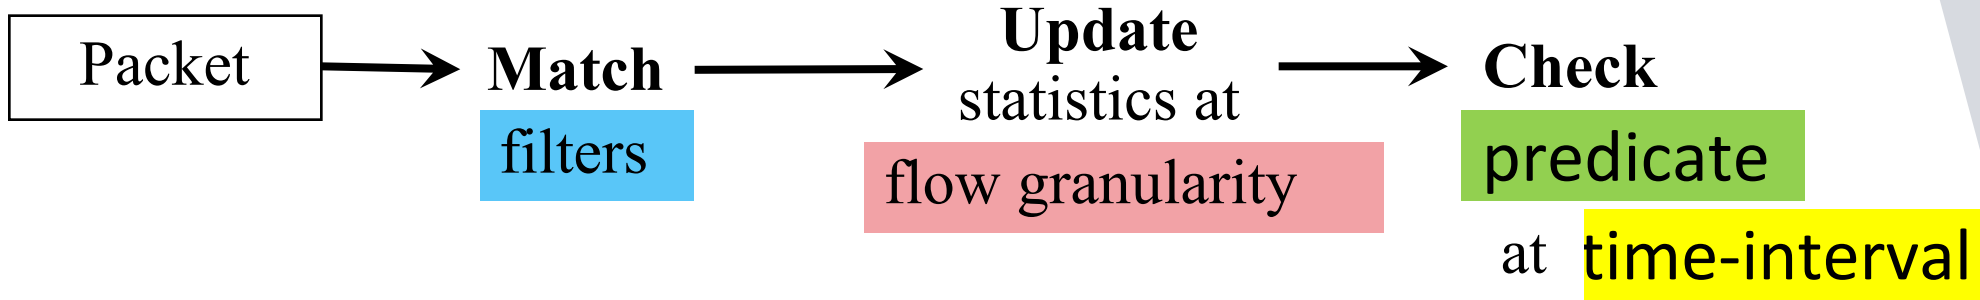


Which of these should be performed

- ❖ *On-path*
- ❖ *Off-path*

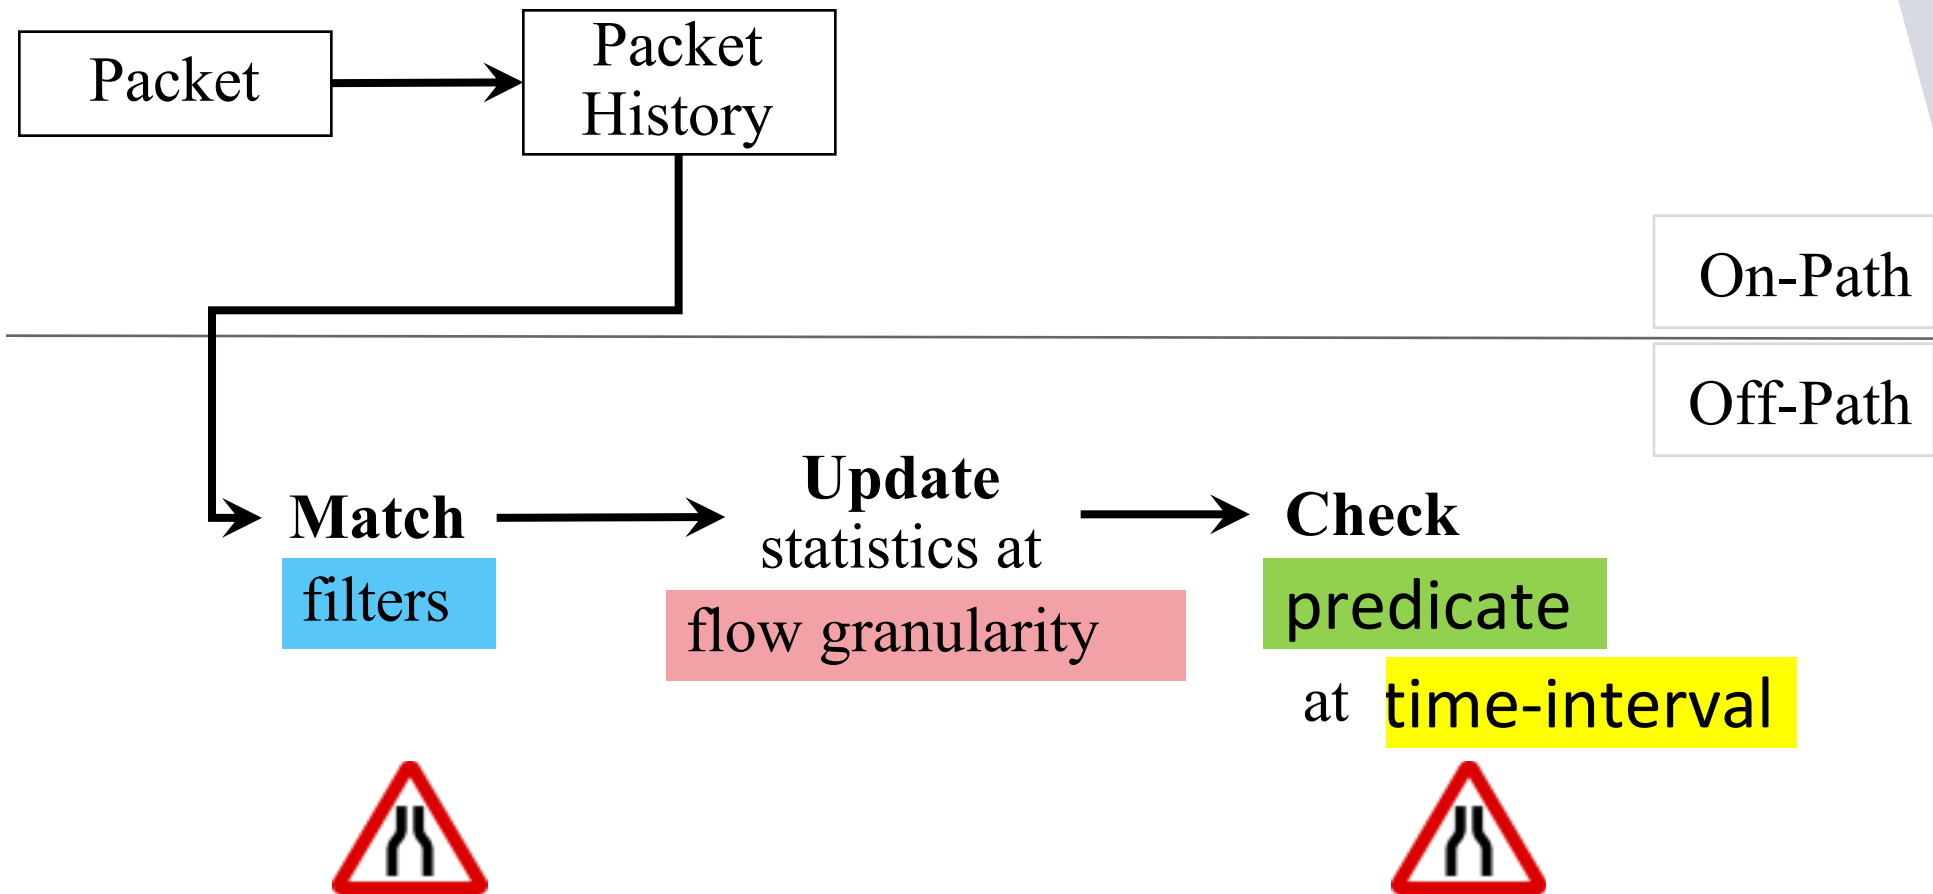


Which operations to do on-path?
❖ 70ns to forward and inspect packet

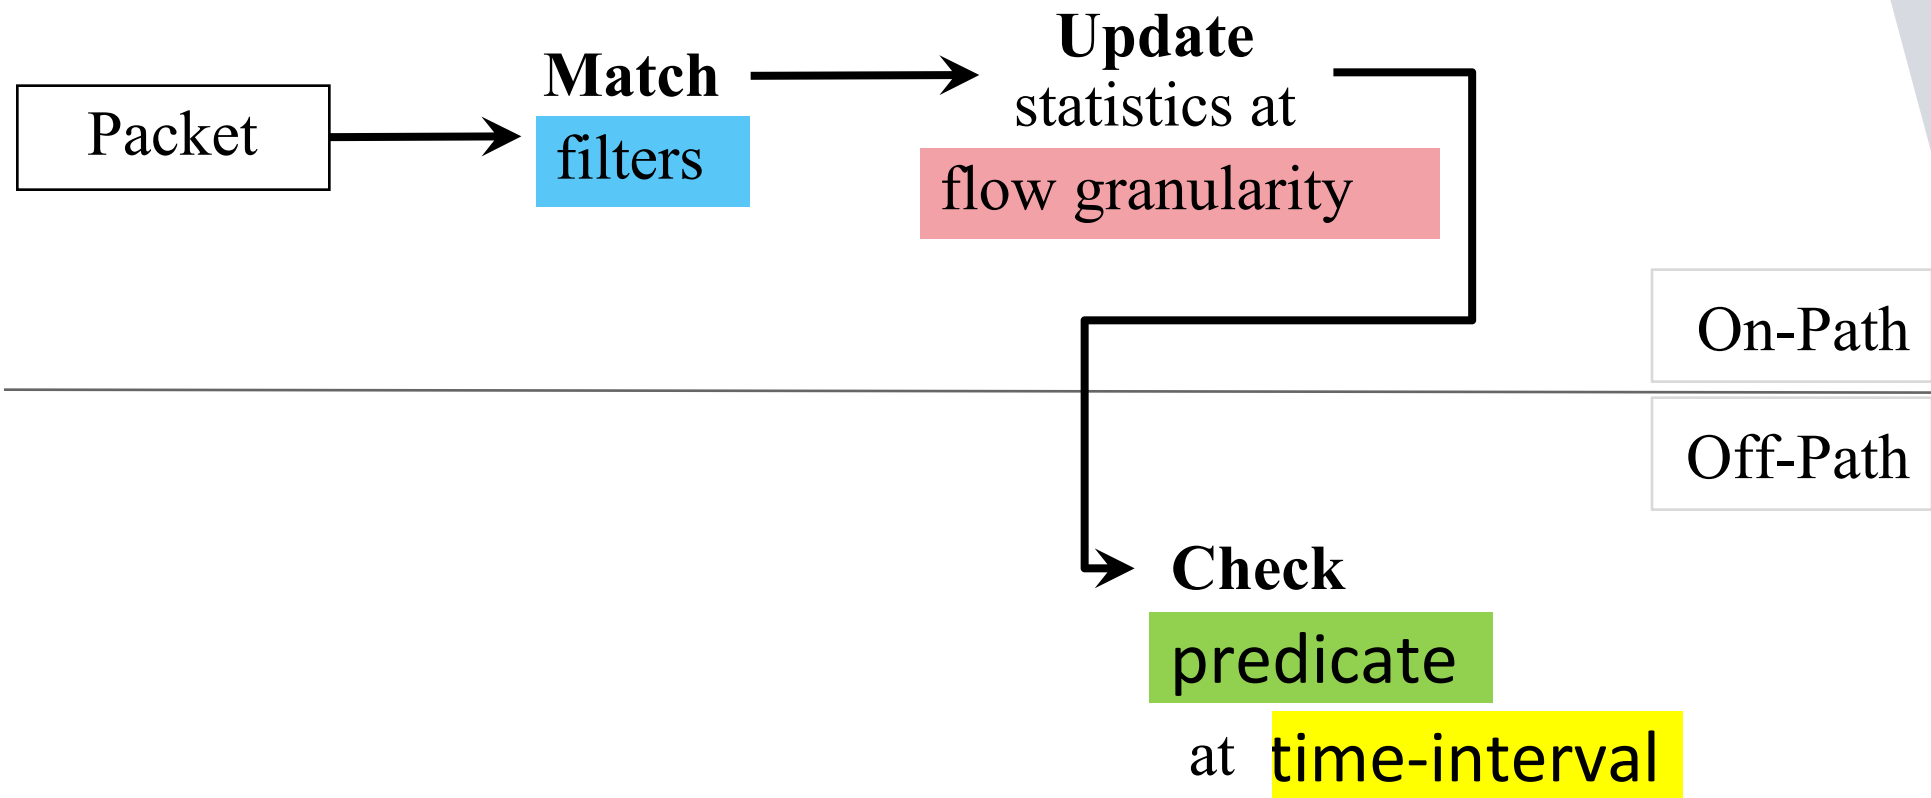


How to schedule off-path operations?

- ❖ Off-path on same core, can delay packets
- ❖ Bound delay to a few μs

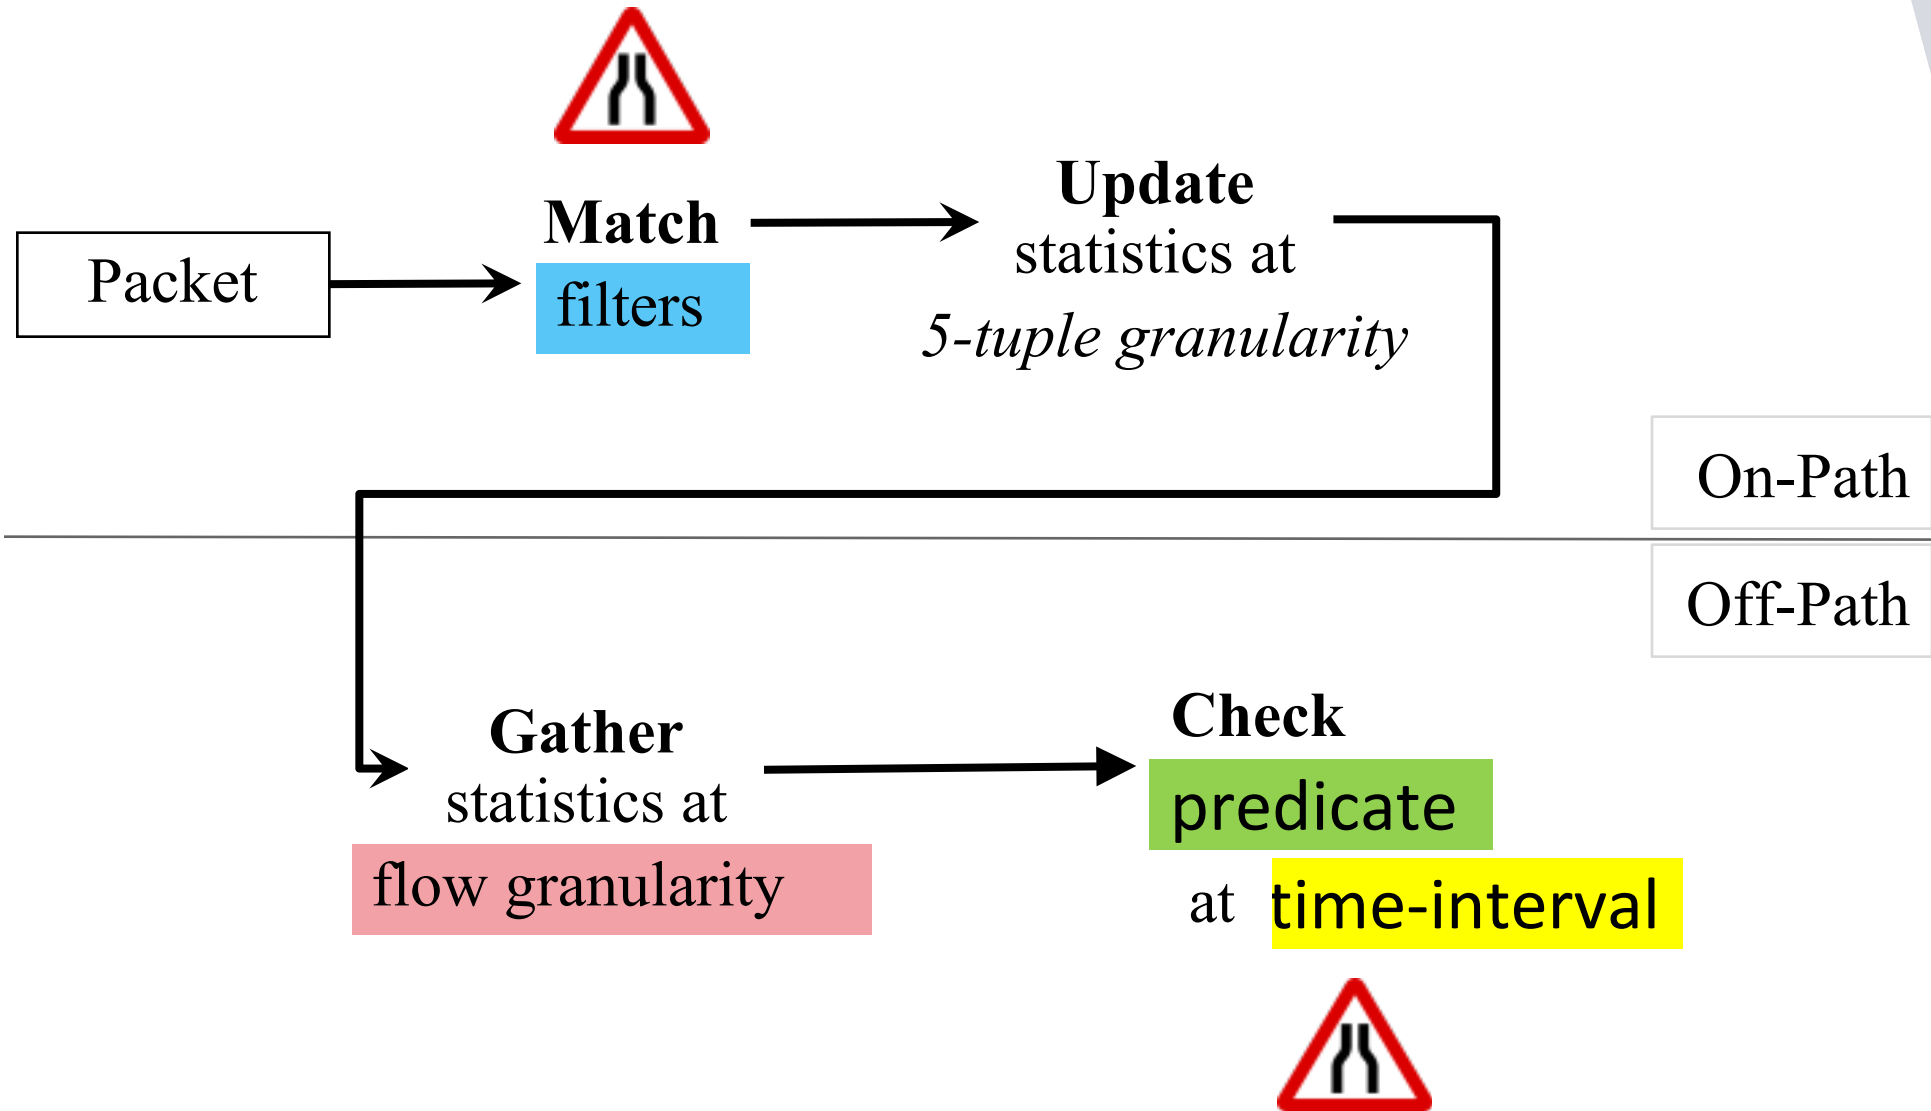



Doesn't scale to large numbers of triggers



Still cannot reach goal

❖ Memory subsystem becomes a bottleneck






On-Path


- ❖ Use tuple-space search for matching
- ❖ Match on first packet, cache match
- ❖ Lay out tables to enable cache prefetch
- ❖ Use TLB huge pages for tables

- ❖ Lazy cleanup of statistics across intervals
- ❖ Lay out tables to enable cache prefetch
- ❖ *Bounded-delay cooperative scheduling*

Off-Path



Bound
delay to
a few μs



Bounded
Delay
Cooperative
Scheduling

Research Questions

What eventing architecture permits programmability and visibility?

How can we achieve precise eventing at fine timescales?

What is the *performance envelope* of such an eventing framework?

Trumpet can monitor thousands of triggers at full packet rate on a 10G NIC

Trumpet is expressive

- ❖ Transient congestion


- ❖ Burst loss

- ❖ Attack onset

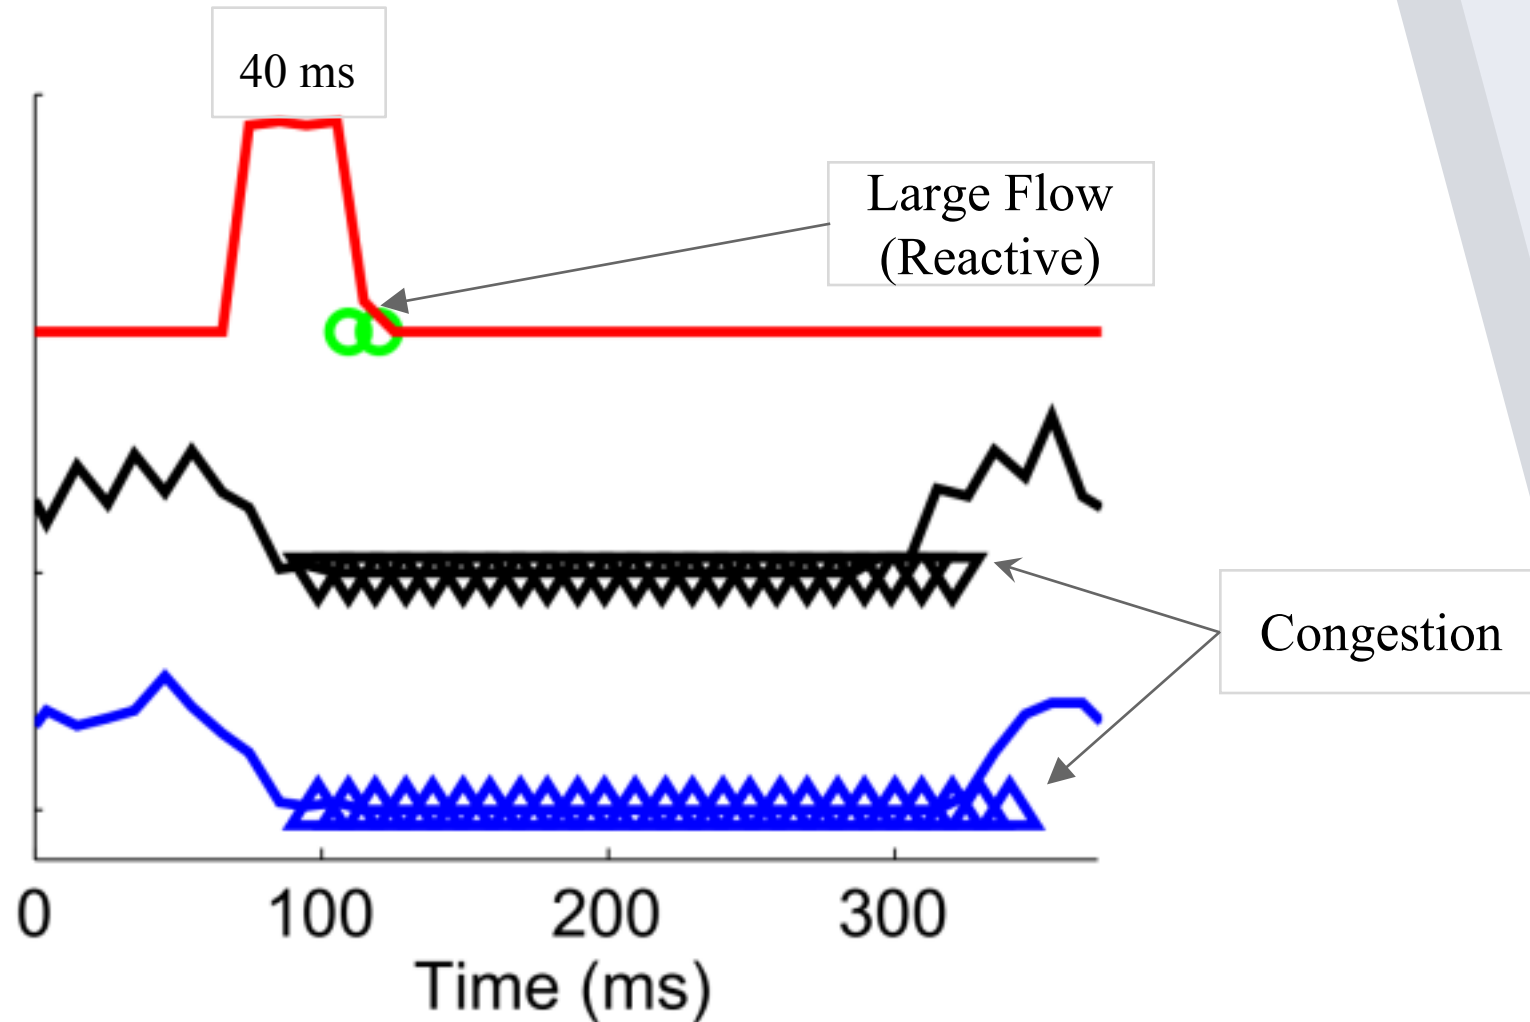
Trumpet scales to thousands of triggers

Trumpet is DoS-Resilient

Detecting Transient Congestion



Trumpet can detect millisecond scale congestion events



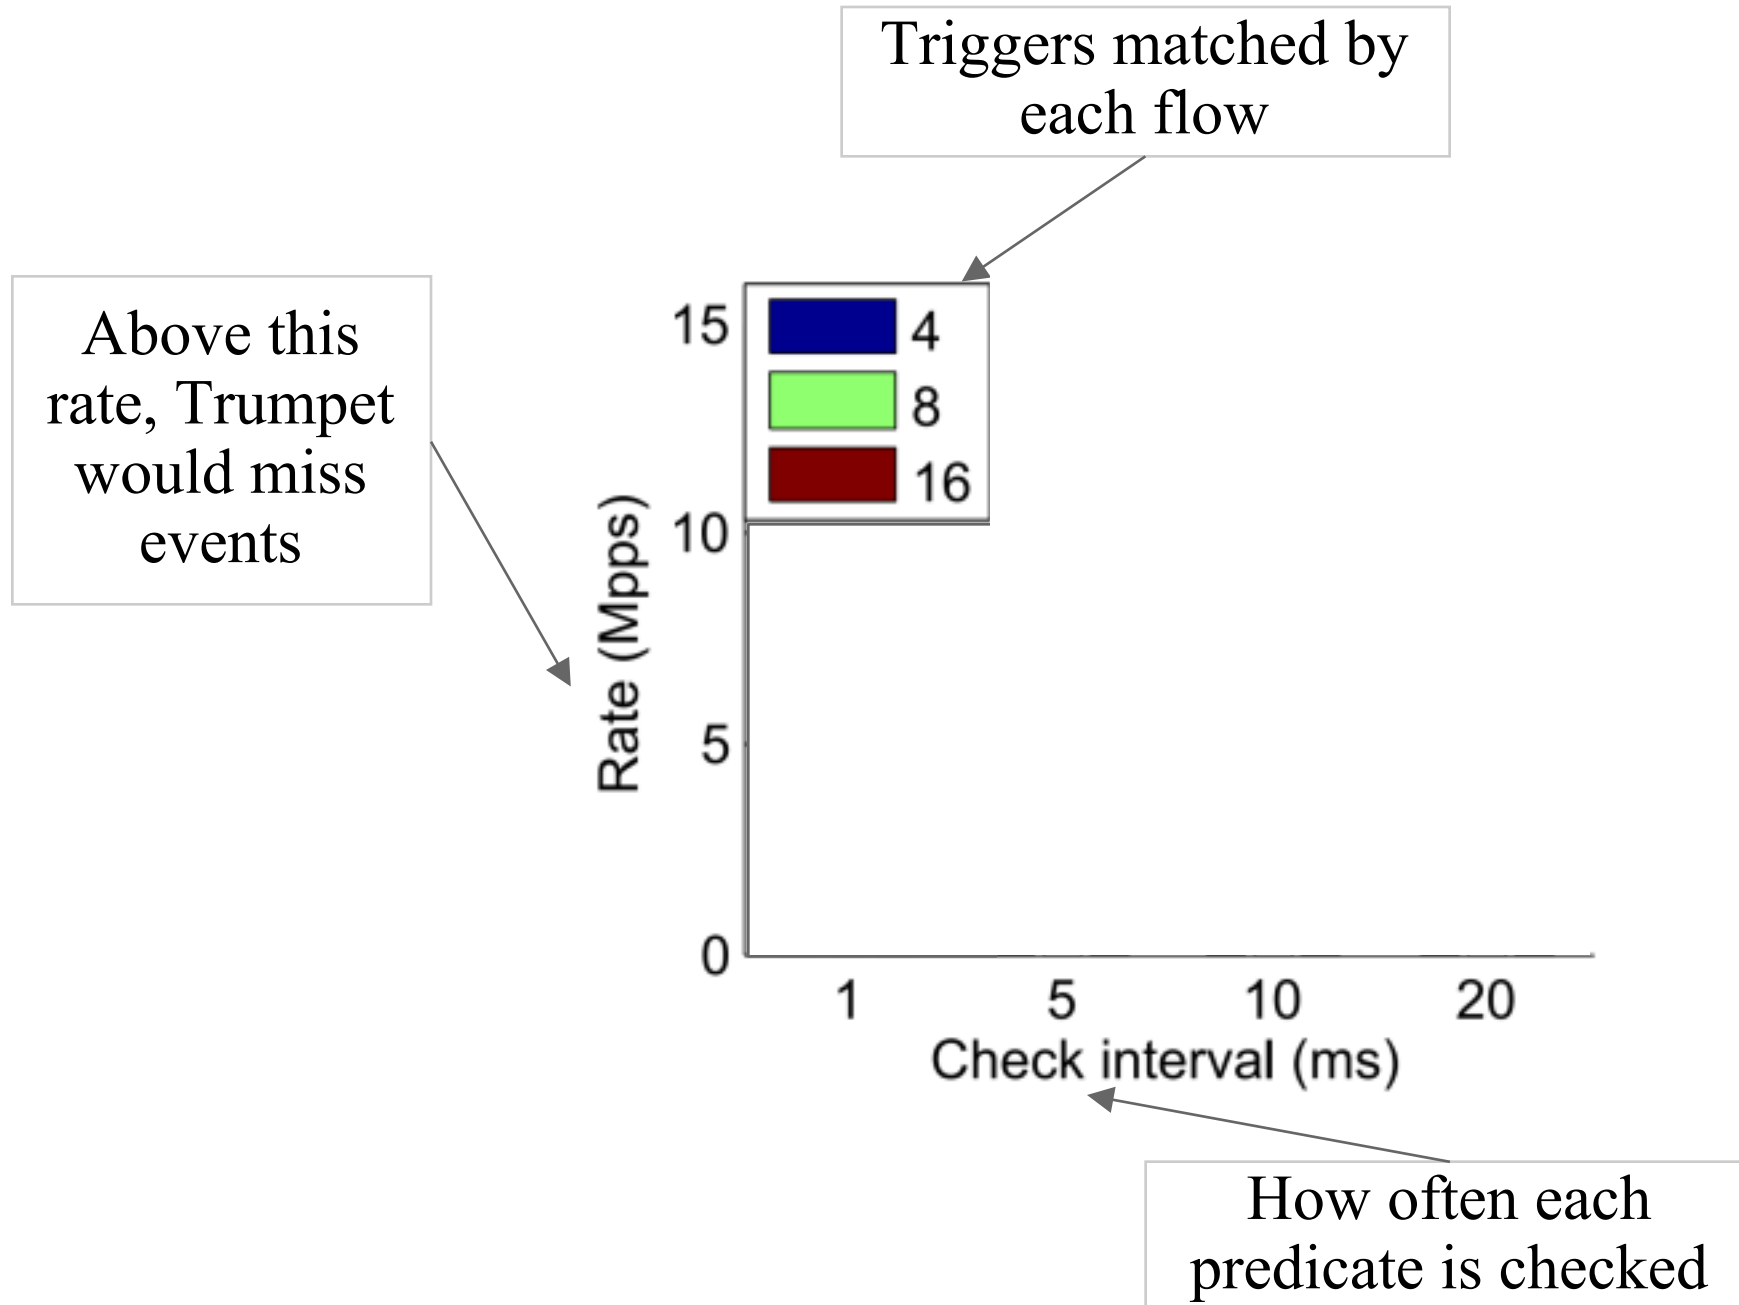
Trumpet can process* **14.8 Mpps**


❖ 64 byte packets at 10G

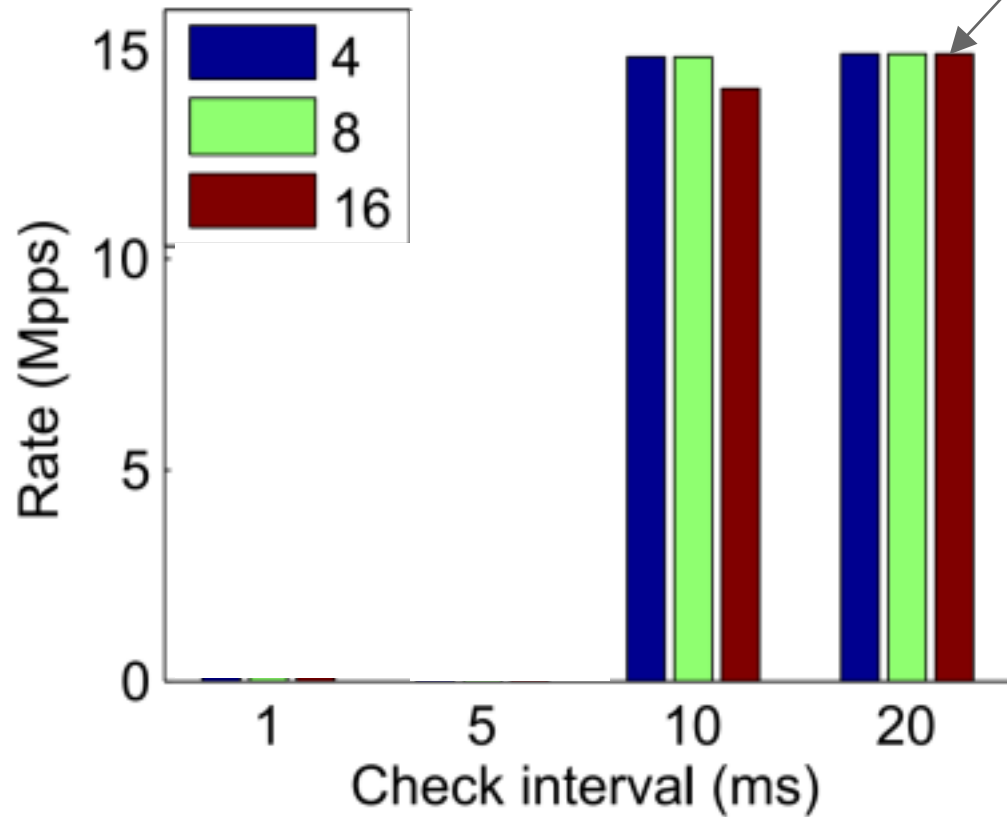
❖ 650 byte packets at 4x10G

... while evaluating **16K triggers** at
10ms granularity

*Xeon ES-2650, 10-core 2.3 Ghz, Intel 82599 10G NIC







Above 10ms, CPU can sustain full packet rate

Need to profile and provision Trumpet deployment

Conclusion

Future datacenters will need fast and precise eventing

- ▶ Trumpet is an expressive system for host-based eventing

Trumpet can process 16K triggers at full packet rate

- ▶ ... without delaying packets by more than 10 μ s

Future work: scale to 40G NICs

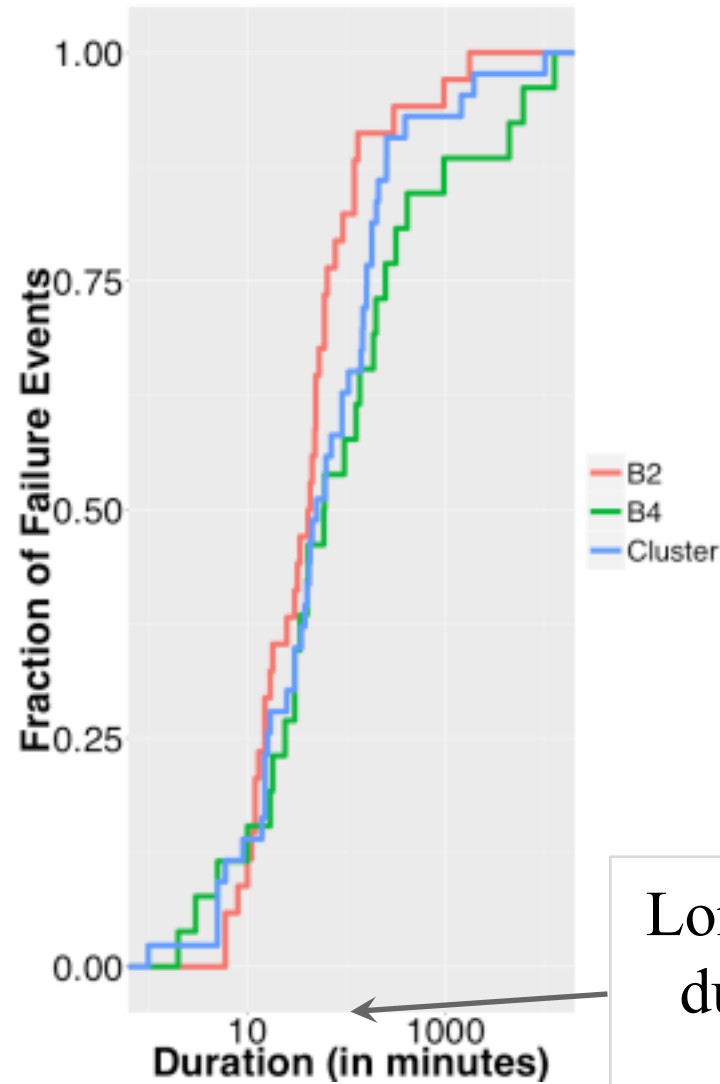
- ▶ ... perhaps with NIC or switch support

<https://github.com/USC-NSL/Trumpet>

Outage budget for **five 9s** availability



24 seconds per month



Long failure durations
due to time to **root-
cause** failures

*Every optimization is necessary**

*Details in the paper