

Programming Network Dataplanes Using P4

Changhoon Kim

P4 Workshop




Agenda

- High-level overview of P4
- Current status of P4
 - Language Spec (v1.0.2)
 - Opensource contributions
- Demo
- Exciting use cases

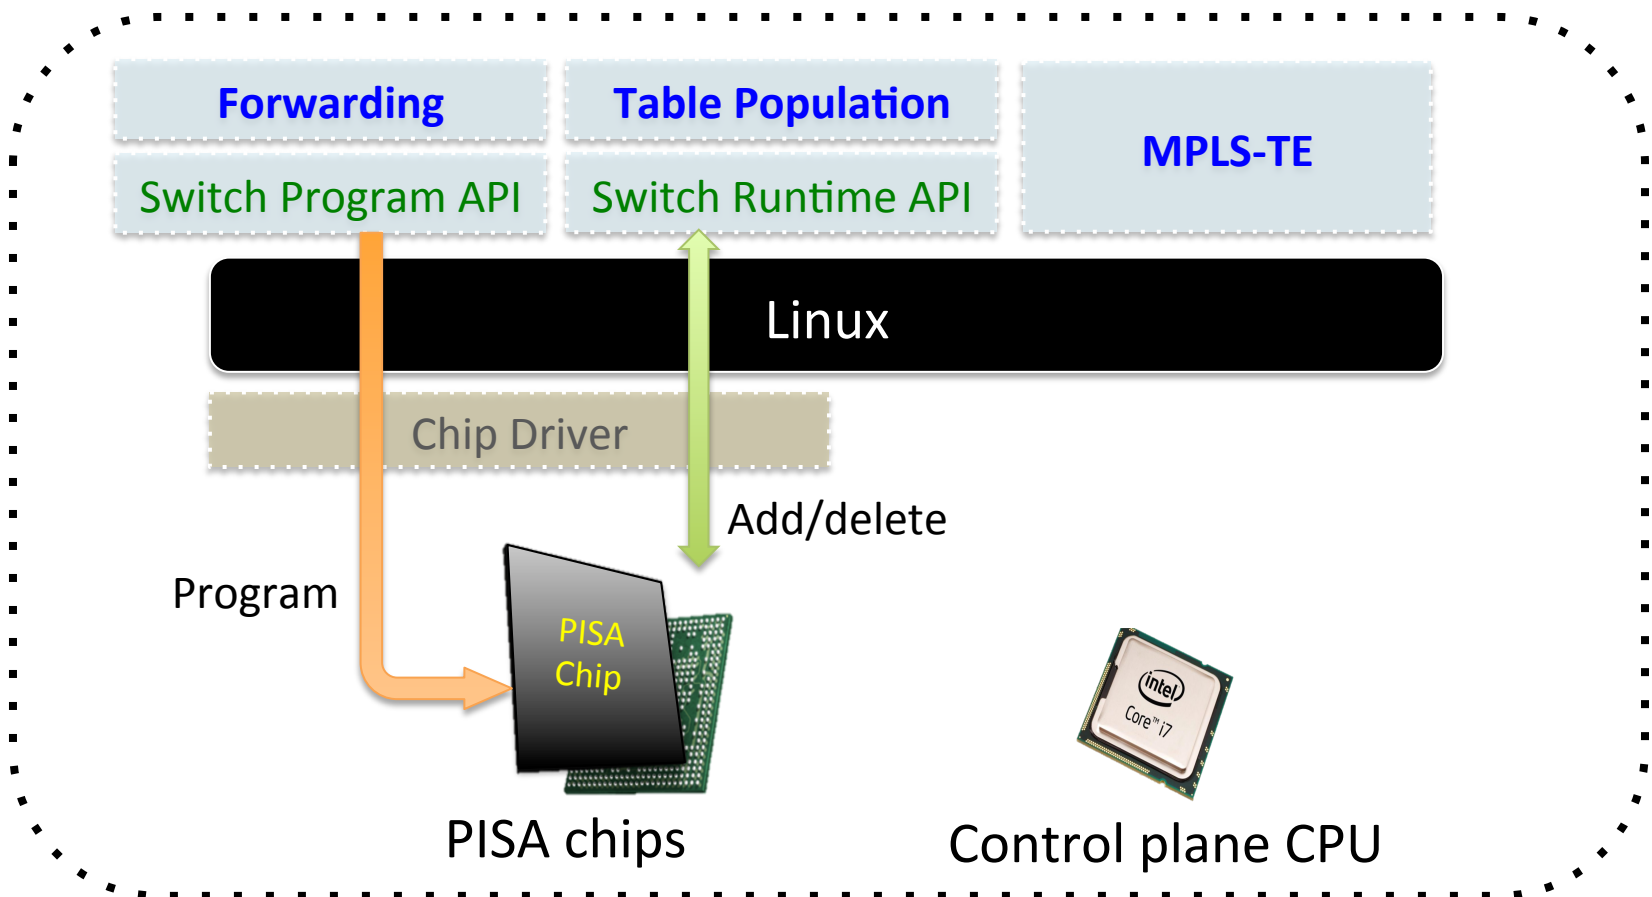
Key players in this new game

- 1. Programmable packet-forwarding solutions
 - S/W switches and GPUs (tens of Gb/s)
 - NPUs and FPGAs (several hundreds of Gb/s)
 - Protocol-independent switching chips (multi Tb/s)
- 2. High-level language
 - P4 (p4.org)
- 3. Compiler and other development tools
 - Common front-end and target-specific back-end compilers
 - Debuggers, profilers, etc.

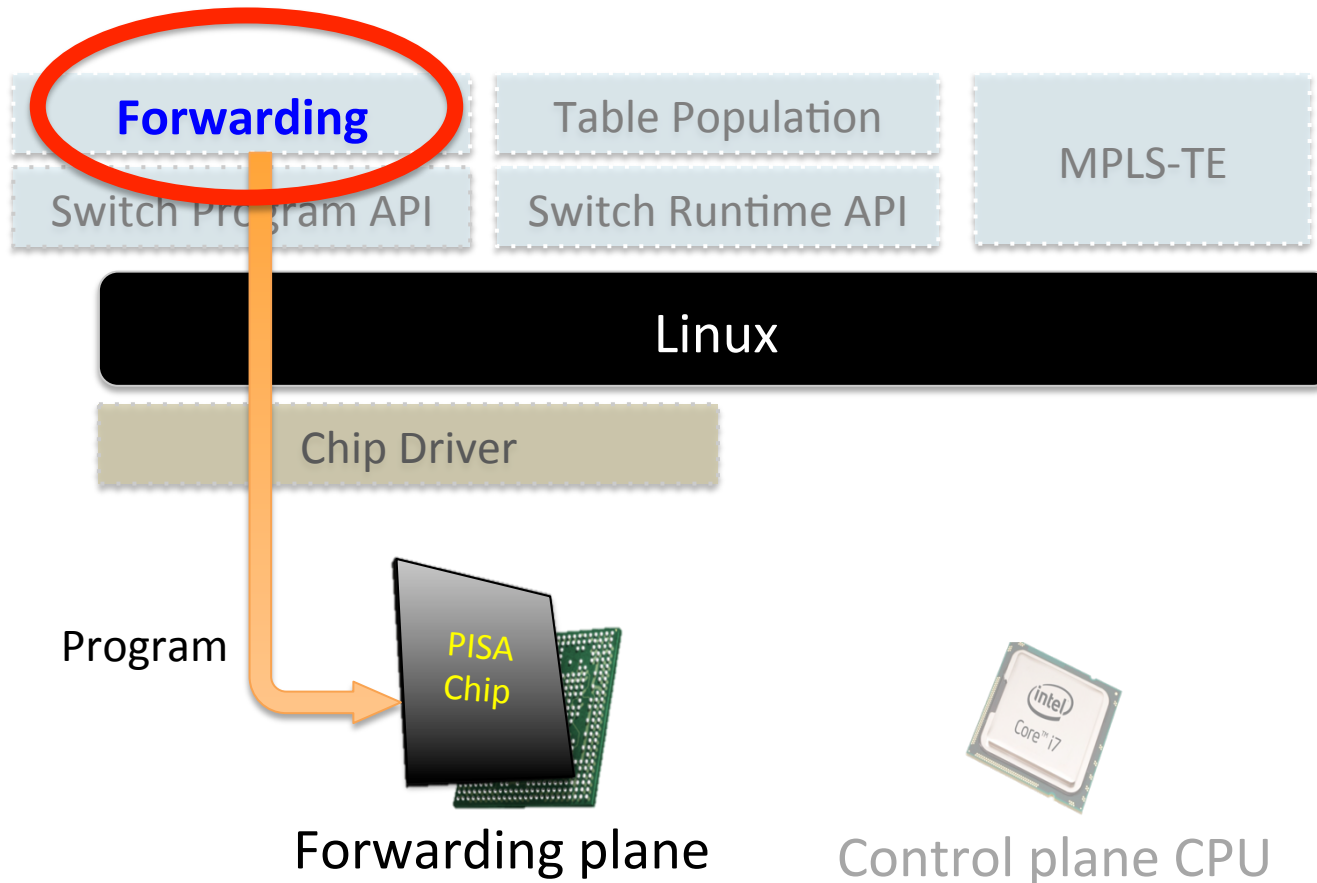
MPLS (source)



My (running) switch



What language to use?



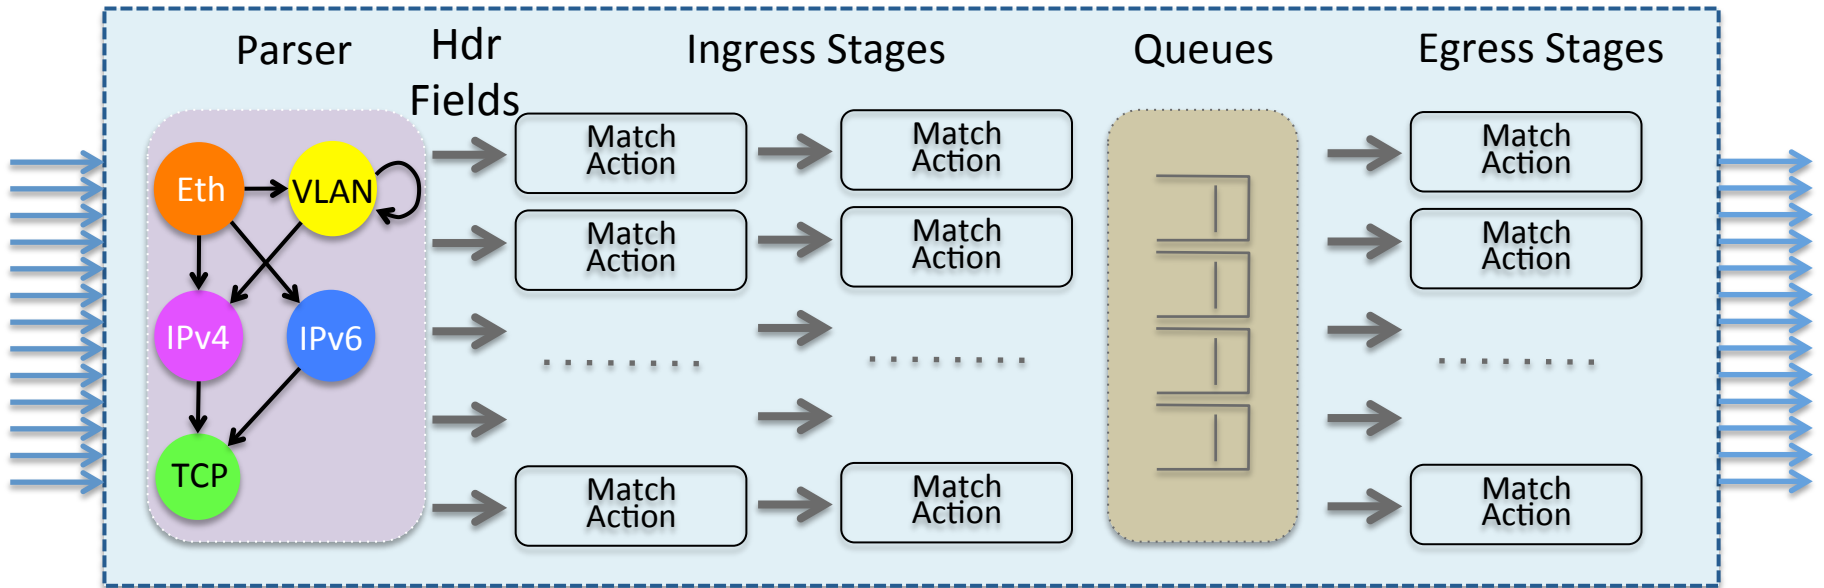
P4: Programming Protocol-independent Packet Processors

Initial contributions from Stanford, Princeton,
Intel, Google, Microsoft, and Barefoot Networks

Goals of P4

- Protocol independence
- Target independence
- Re-configurability in the field
- (And, intuitive abstractions)

Abstract Forwarding Model



P4 Language Components

**Parser
Program**

State-machine;
Field extraction

**Match Tables
+ Actions**

Control Flow

Table lookup and update;
Field manipulation;
Control flow

**Assembly
("deparser")
Program**

Field assembly

No: memory (pointers), loops, recursion, floating point


P4 Code Examples

1. Header Fields and Parsing
2. Match + Action Tables
3. Control flow

Header Fields and Parsing

```
header_type ethernet_t {  
    fields {  
        dstAddr : 48;  
        srcAddr : 48;  
        etherType : 16;  
    }  
}
```


```
parser parse_ethernet {  
    extract(ethernet);  
    return select(latest.etherType) {  
        0x8100 : parse_vlan;  
        0x800 : parse_ipv4;  
        0x86DD : parse_ipv6;  
    }  
}
```



Match

```
table ipv4_lpm
{
    reads {
        ipv4.dstAddr : lpm;
    }
    actions {
        set_next_hop;
        drop;
    }
}
```

Lookup key



ipv4.dstAddr	action
0.*	drop
10.0.0.*	set_next_hop
224.*	drop
192.168.*	drop
10.0.1.*	set_next_hop


Actions

ipv4.dstAddr	action
0.*	drop
10.0.0.*	set_next_hop
224.*	drop
192.168.*	drop
10.0.1.*	set_next_hop

nhop_ipv4_addr	port
10.0.0.10	1
10.0.1.10	2

```
action set_next_hop(nhop_ipv4_addr, port)
{
    modify_field(metadata.nhop_ipv4_addr, nhop_ipv4_addr);
    modify_field(standard_metadata.egress_port, port);
    add_to_field(ipv4.ttl, -1);
}
```



Control Flow



```
control ingress
```

```
{  
    apply(port);  
    if (valid(vlan_tag[0])) {  
        apply(port_vlan);  
    }  
    apply (bridge_domain);  
    if (valid(mpls_bos)) {  
        apply(mpls_label);  
    }  
    retrieve_tunnel_vni();  
    if (valid(vxlan) or valid(genv) or valid(nvgre))  
    {  
        apply(dest_vtep);  
        apply(src_vtep);  
    }  
}
```

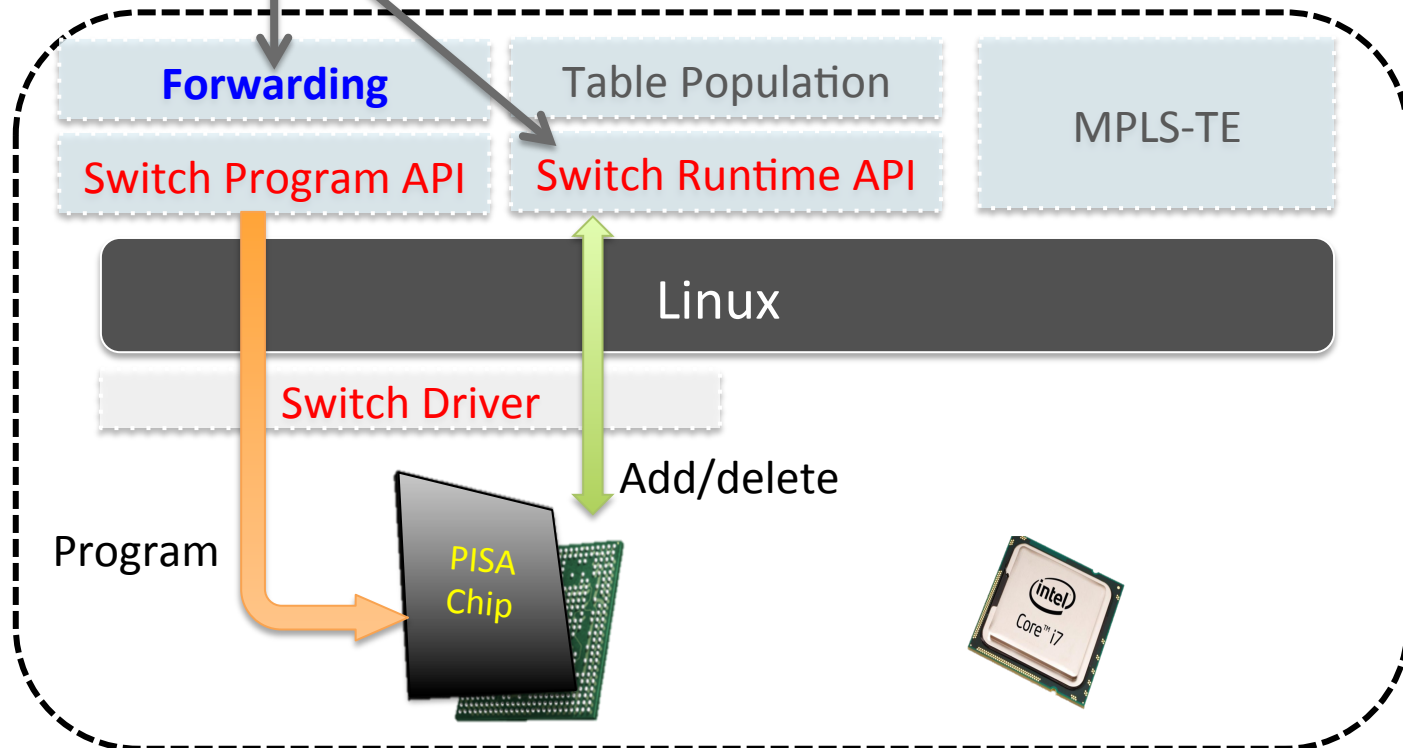
MPLS (source)



Compiler



My (running) switch



Switch.p4: A datacenter switch in P4

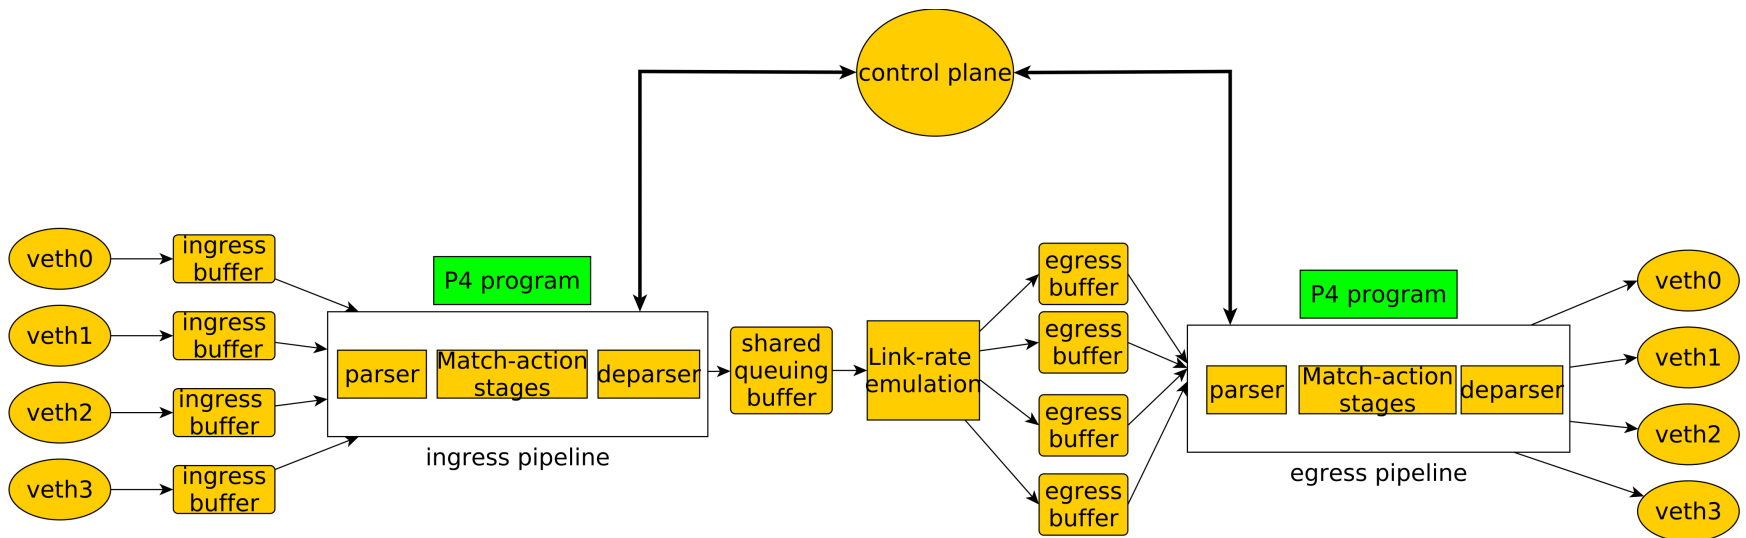
- Can we express the forwarding plane of a typical DC switch in P4?
- Feature set
 - Basic L2 switching: MAC learning, VLAN, flooding, and STP
 - Basic L3 routing: IPv4 and IPv4 routing with VRF
 - LAG
 - ECMP
 - Tunneling: VXLAN and NVGRE (L2/L3 Gateway), Geneve, and GRE
 - Basic ACL: MAC and IP ACLs
 - Unicast RPF
 - MPLS: LER, LSR, IPVPN, VPLS, and L2VPN
- More features coming soon

Opensource P4 Dev Environment

- Sample P4 programs
- Compiler
- S/W P4 switch
- Test framework

S/W P4 Switch

- Stand-alone test framework
- Network-level test framework
 - Pluggable into Mininet through virtual interfaces



Original P4 Paper → Spec v1.0.2

- Primitive actions
 - clone() variants
 - drop()
 - generate_digest()
 - modify_field_with_hash()
- Action profile
- Field list
- Header stacks
- Counters, meters, and registers

New Language Constructs Enabled More Features

- ECMP
- MAC learning
- ACLs
- Mirroring
- MPLS
- Statistics
- Various hash-based algorithms

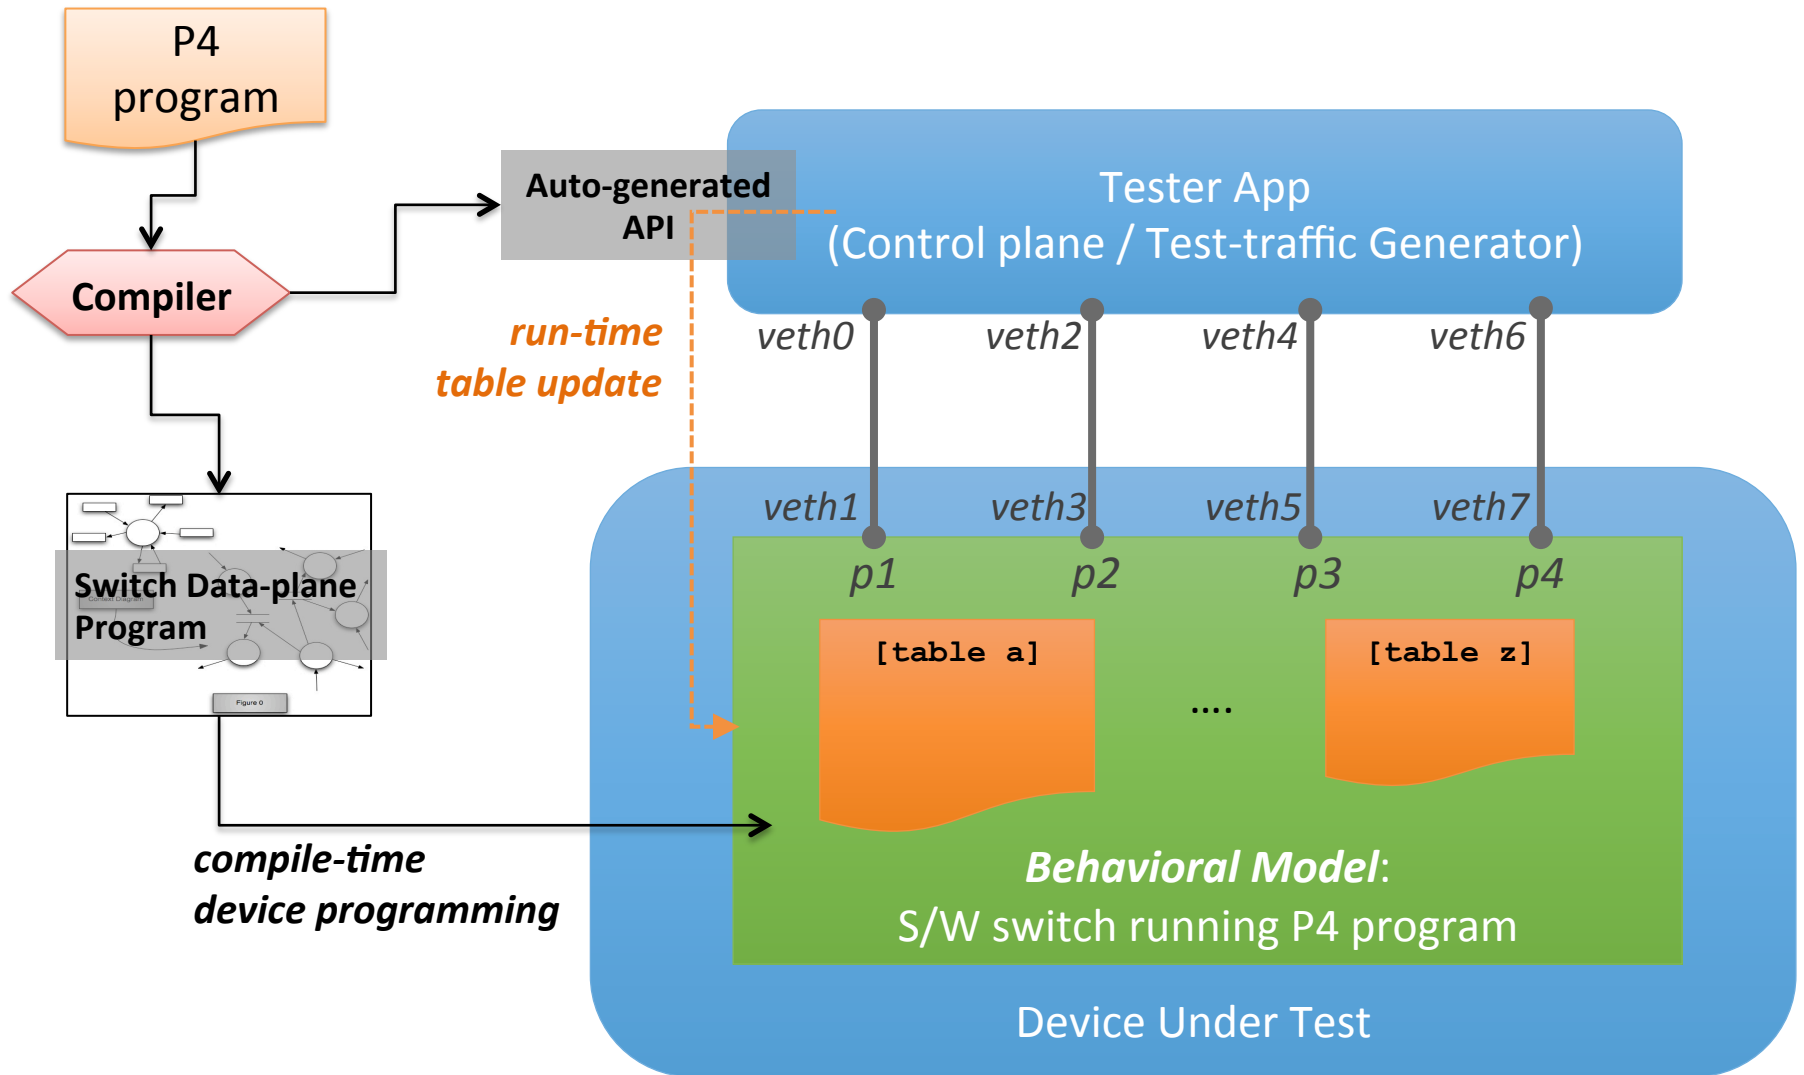
Avenues Where More Contributions Are Needed

- Short term
 - Constructs for TLV-style header parsing
 - Register manipulation
- Long term
 - Modeling buffers and schedulers
 - Buffer management schemes
 - Carving up buffers into queues and application pools
 - Packet scheduling schemes
 - PQ, WFQ, HFQ, etc.

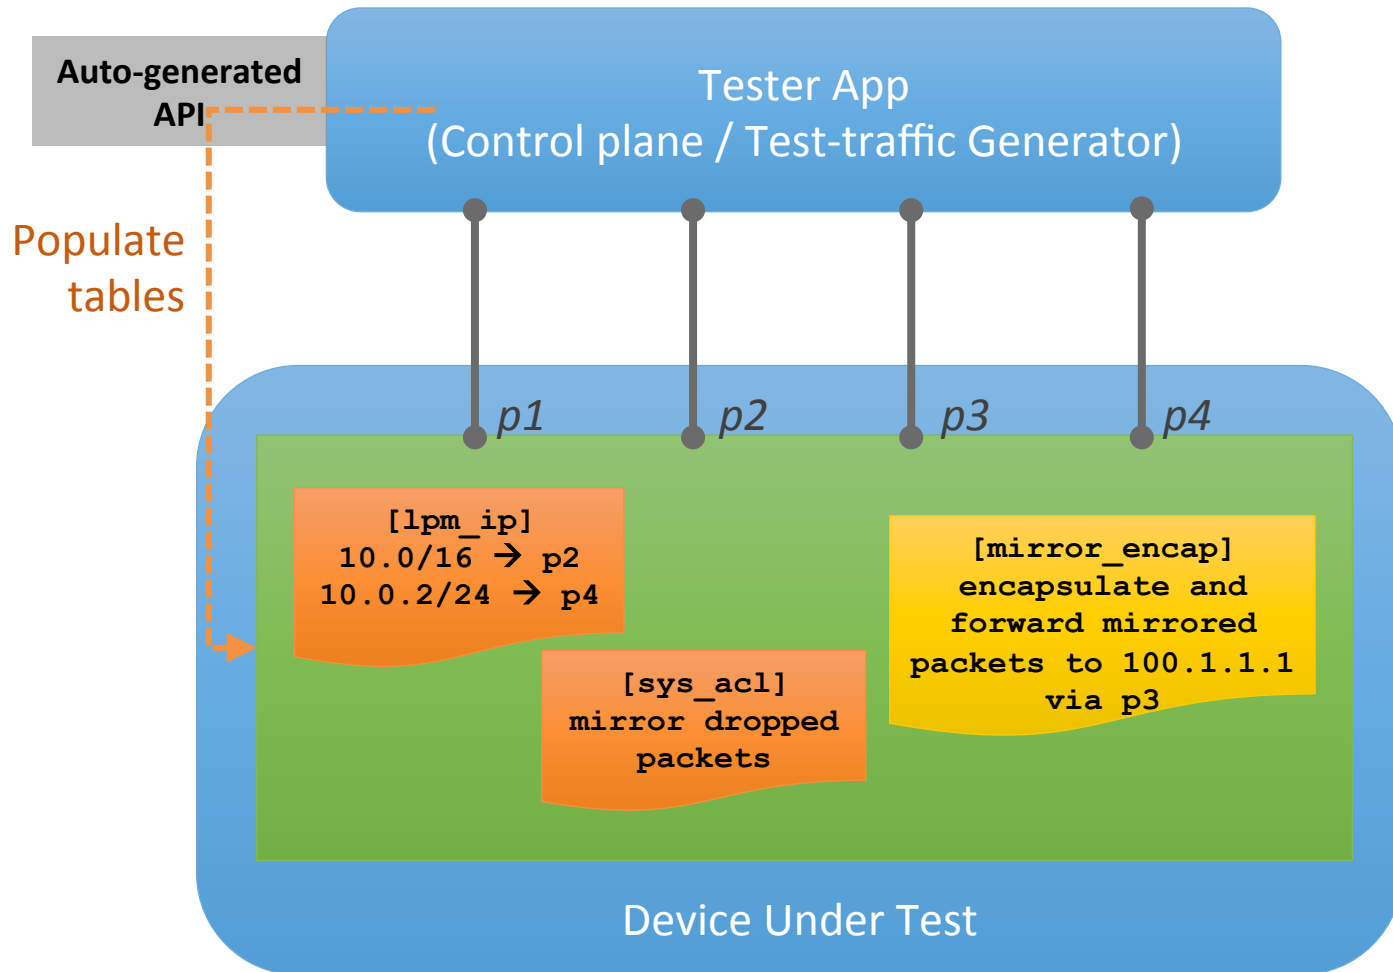
DEMO!

Reproducible via opensource dev tools and
reference P4 programs

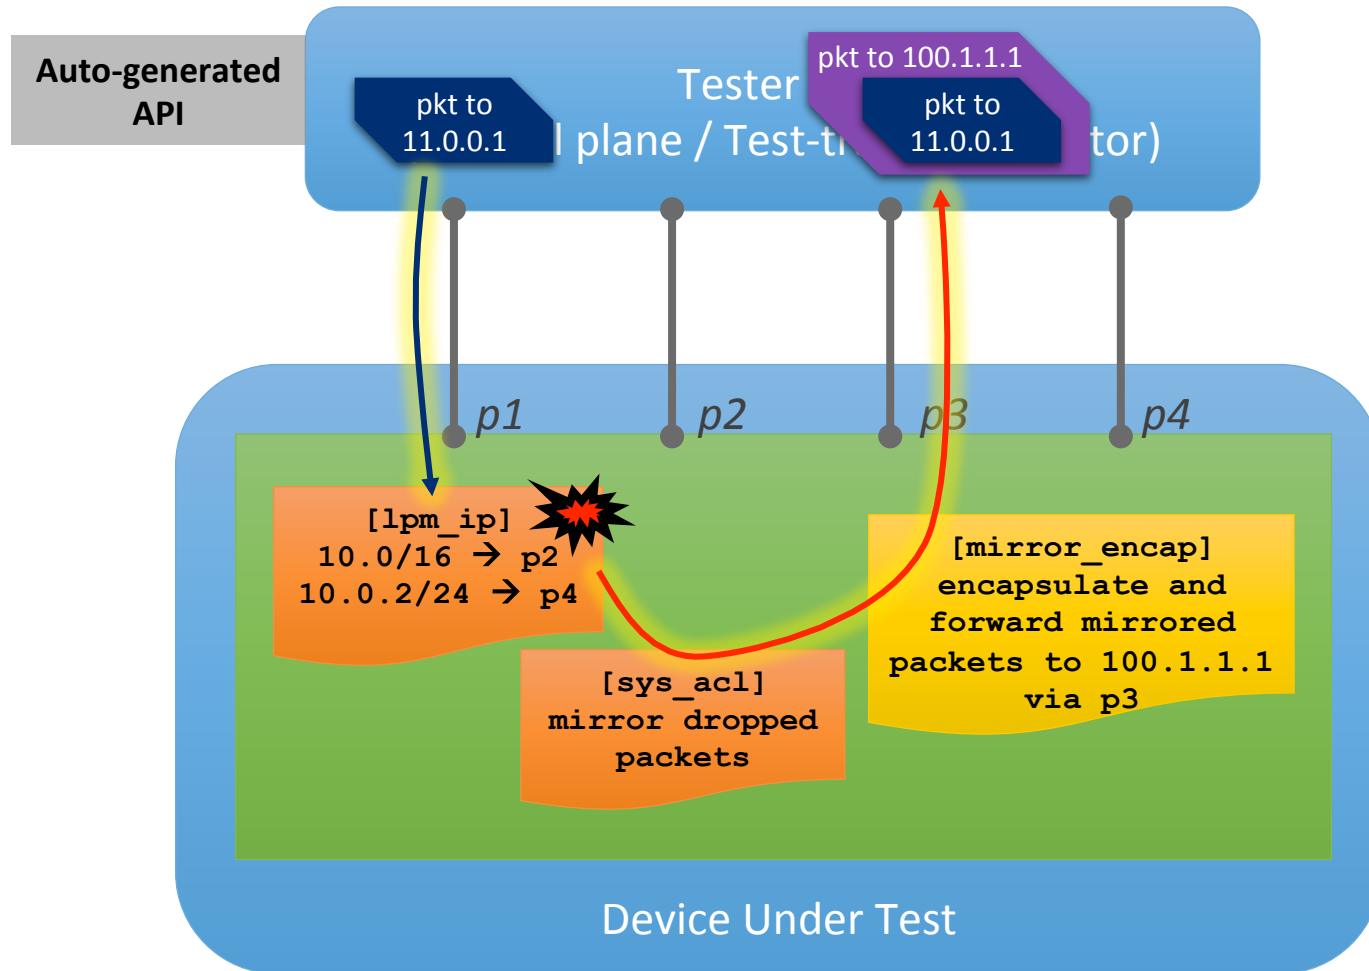
Demo Setup



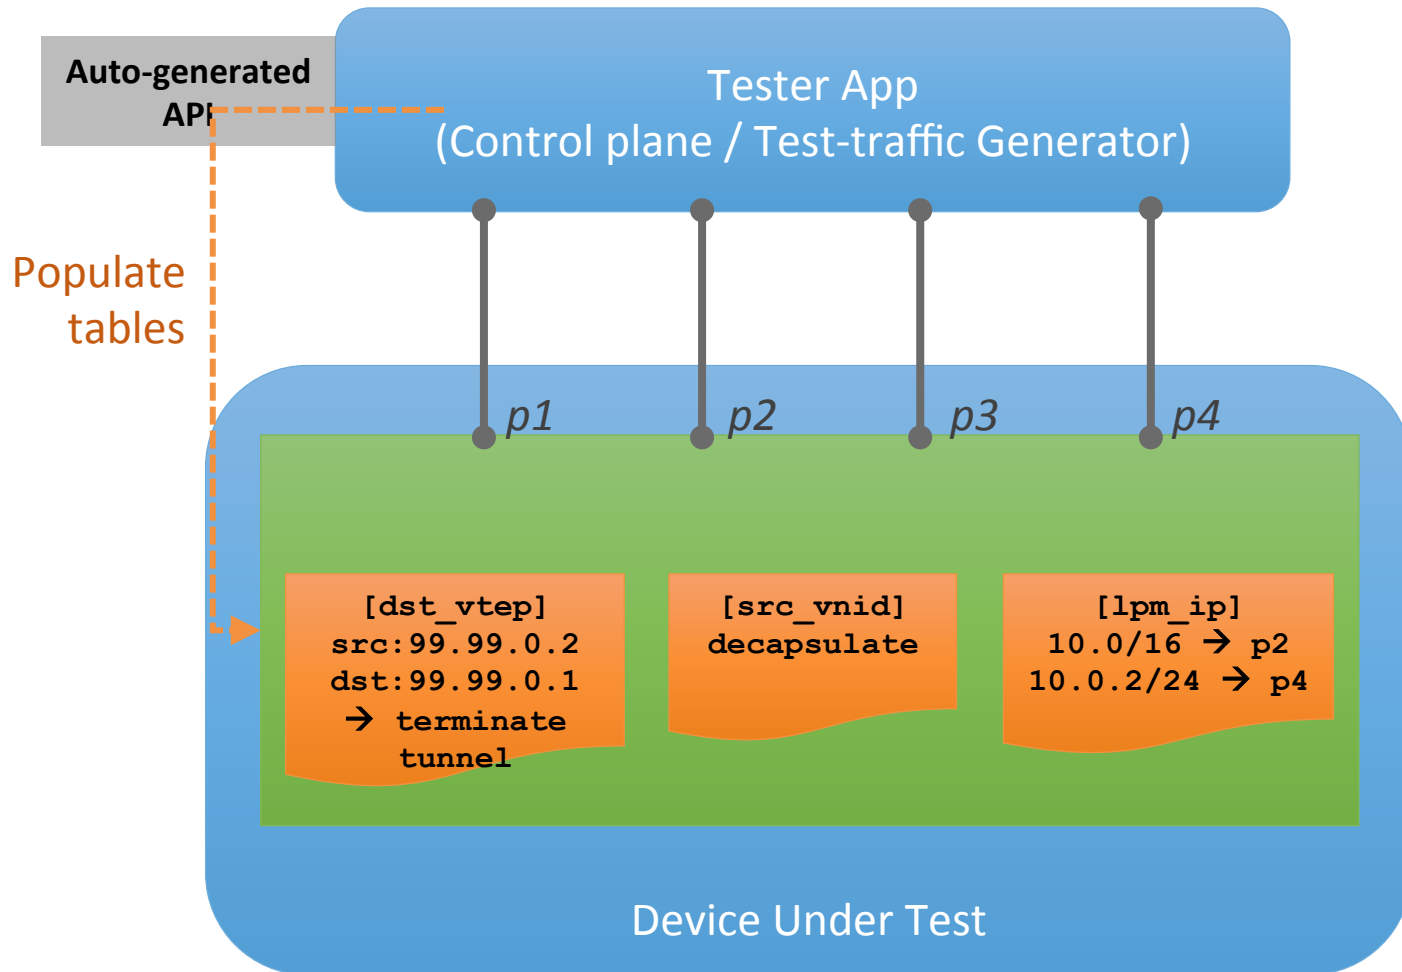
Scenario 1: Mirroring Dropped Packets



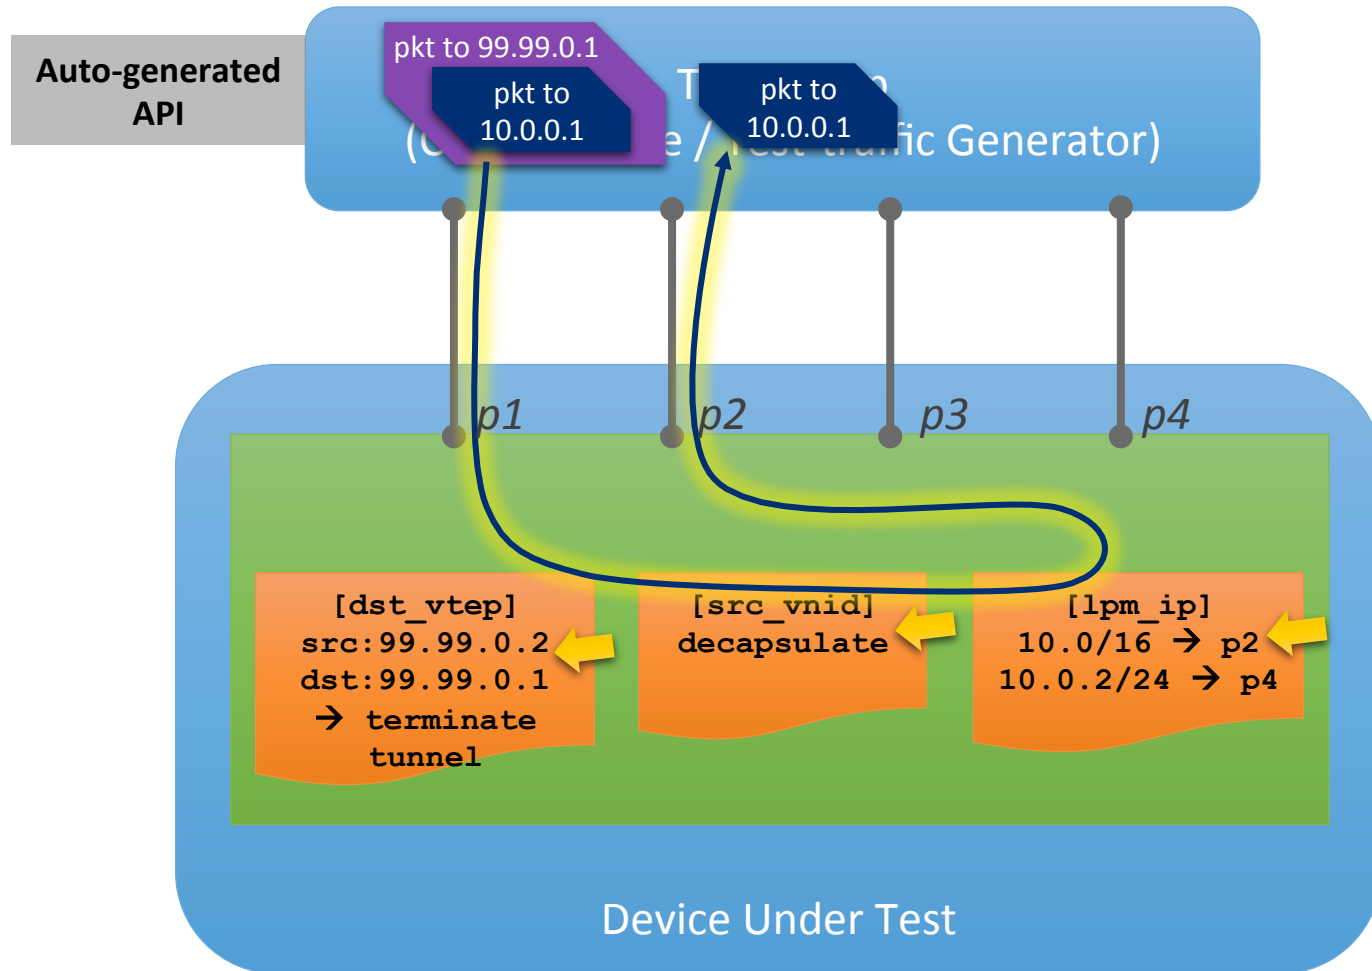
Scenario 1: Mirroring Dropped Packets



Scenario 2: VXLAN Routing



Scenario 2: VXLAN Routing



Exciting Applications

- In-band Network Telemetry (INT)
 - Highly user for debugging and analysis
- Tunnel splicing
 - P2V gateways
- L4 load-balancing
- Utilization-aware fabric load-balancing

INT Specification

- INT v1.5 is available to P4.org members
 - Includes reference P4 code
- Participants
 - VMware, Barefoot, and a few other companies
- We would appreciate feedback from the P4 community

Summary

- Common industry-wide forwarding language will be immensely helpful
- P4 looks very promising
- How can you get involved?
 - Join P4.org
 - Assign engineers to get familiar with P4
 - Start playing tools, language, and sample P4 code
 - Contribute back to P4.org

Full-day tutorial at SIGCOMM in Aug at London!

